

# Game Theory and Security

## Human Behavior Modeling & Learning

---

Fei Fang  
Harvard University  
Carnegie Mellon University

[feifang@cmu.edu](mailto:feifang@cmu.edu)

# Wildlife Population Threatened by Poaching



Data SIO, NOAA, U.S. Navy, NGA, GEBCO  
Image Landsat  
Image IBCAO

Google earth

# Wildlife Population Threatened by Poaching



Today

≈ 3,200



100 years ago

≈ 60,000



# Wildlife Population Threatened by Poaching

- ▶ Get the most out of the patrols
  - ▶ Game theory + machine learning



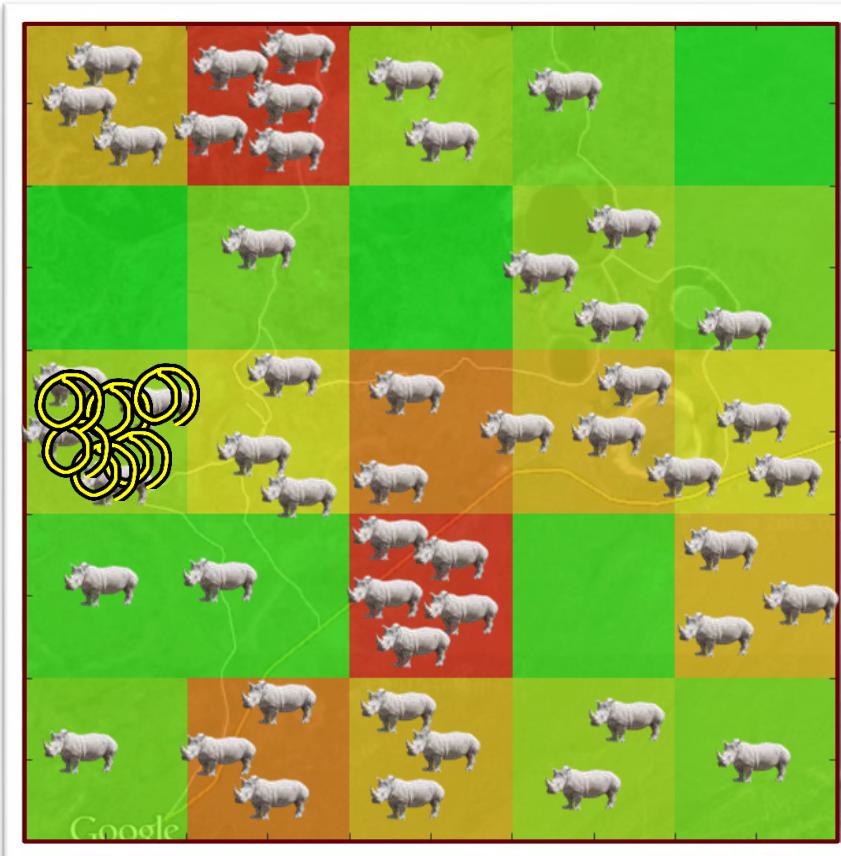
# Challenges in Wildlife Conservation

- ▶ Frequent and repeated attacks
  - ▶ Not one-shot
- ▶ Attacker decision making
  - ▶ Limited surveillance / Less effort / Boundedly rational
- ▶ Real-world data
  - ▶ Sparse / Incomplete / Uncertainty / Noise
- ▶ Real-world deployment
  - ▶ Practical constraints
  - ▶ Field test



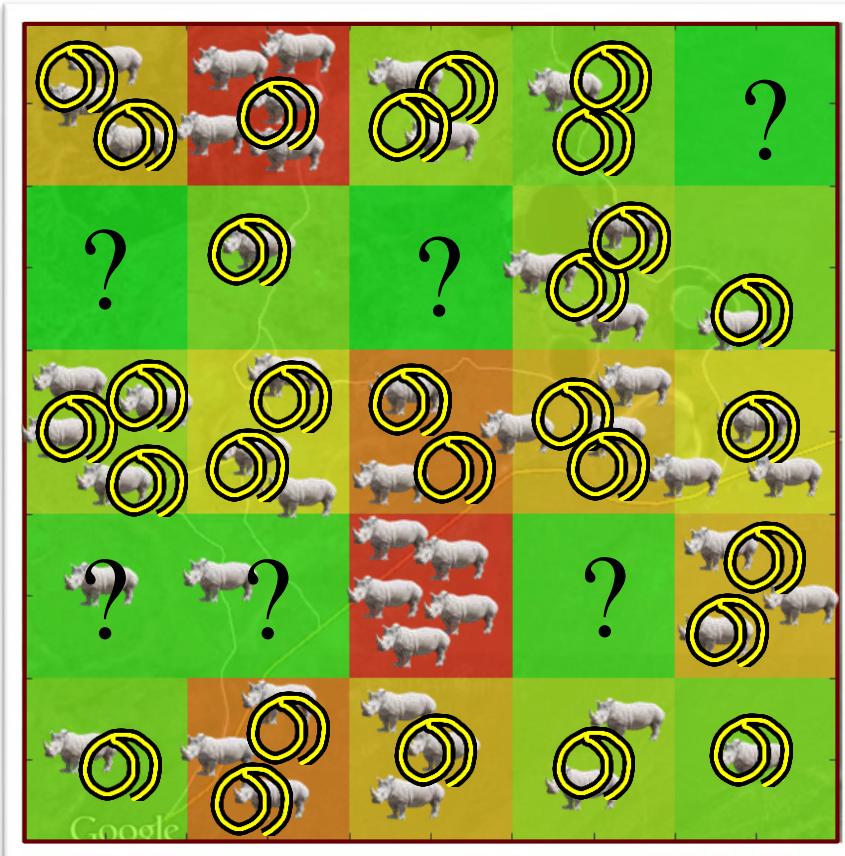
# Challenges in Wildlife Conservation

- ▶ Perfectly rational (Maximize expected utility)? No!



# Challenges in Wildlife Conservation

## ► Real-world data



# Human Behavior Modeling & Learning

- ▶ Uncertainty and Bias Based Models
  - ▶ Prospect Theory [Kahneman and Tversky, 1979]
  - ▶ Anchoring bias and epsilon-bounded rationality [Pita et al, 2010]
  - ▶ Attacker aims to reduce the defender's utility [Pita et al, 2012]
- ▶ Quantal Response Based Models
  - ▶ Quantal Response [McKelvey and Palfrey, 1995]
  - ▶ Subjective Utility Quantal Response [Nguyen et al, 2013]
- ▶ Latest Models
  - ▶ Incorporating delayed observation [Fang et al, 2015]
  - ▶ Bounded rationality in repeated games [Kar et al, 2015]
  - ▶ Two-layered model [Nguyen et al, 2016]
  - ▶ Decision tree-based model [Kar & Ford et al, 2017]
- ▶ PAWS

# Human Behavior Modeling & Learning

- ▶ **Uncertainty and Bias Based Models**
  - ▶ Prospect Theory [Kahneman and Tversky, 1979]
  - ▶ Anchoring bias and epsilon-bounded rationality [Pita et al, 2010]
  - ▶ Attacker aims to reduce the defender's utility [Pita et al, 2012]
- ▶ **Quantal Response Based Models**
  - ▶ Quantal Response [McKelvey and Palfrey, 1995]
  - ▶ Subjective Utility Quantal Response [Nguyen et al, 2013]
- ▶ **Latest Models**
  - ▶ Incorporating delayed observation [Fang et al, 2015]
  - ▶ Bounded rationality in repeated games [Kar et al, 2015]
  - ▶ Two-layered model [Nguyen et al, 2016]
  - ▶ Decision tree-based model [Kar & Ford et al, 2017]
- ▶ **PAWS**

# PT: Prospect Theory

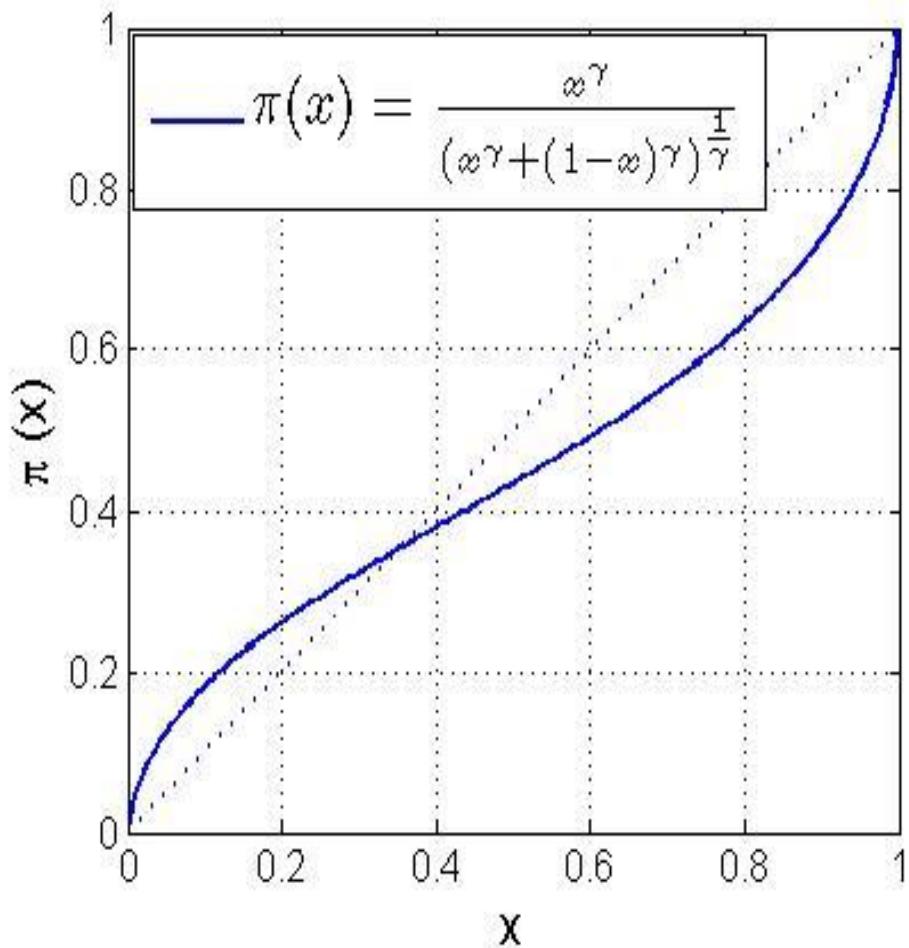
- ▶ Model human decision making under uncertainty
- ▶ Maximize the ‘prospect’ [Kahneman and Tversky, 1979]

$$\text{prospect} = \sum_{i \in \text{AllOutcomes}} \pi(x_i) \cdot V(C_i)$$

- ▶  $\pi(\cdot)$ : weighting function
- ▶  $V(\cdot)$ : value function
- ▶ Defender: choose a strategy that maximizes DefEU when attacker best responds to the expected prospect (instead of AttEU)

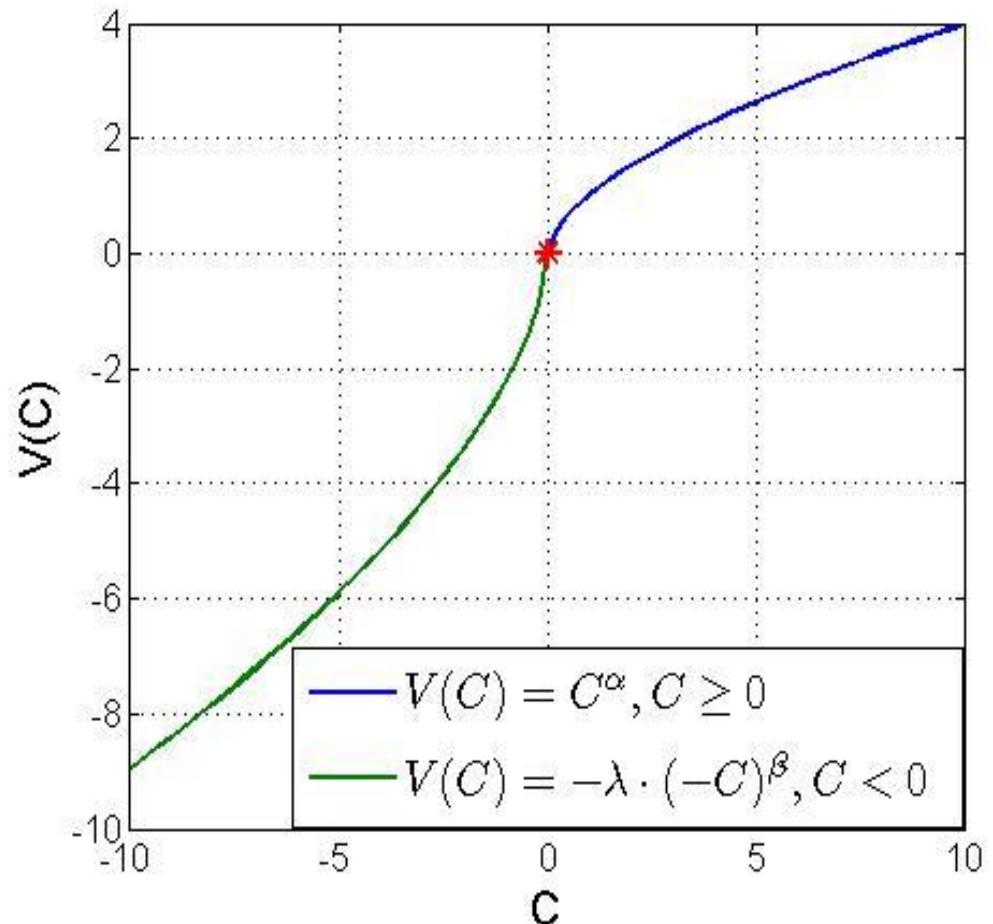
# PT: Prospect Theory

- ▶ Empirical Weighting Function
- ▶ Slope gets steeper as  $x$  gets closer to 0 and 1
- ▶ Not consistent with probability definition
  - $\pi(x) + \pi(1-x) < 1$
- ▶ Empirical value:  
 $\gamma=0.64$  ( $0 < \gamma < 1$ )



# PT: Prospect Theory

- ▶ Empirical Value Function
- ▶ Risk averse regarding gain
- ▶ Risk seeking regarding loss
- ▶ Empirical value:  
 $\alpha=\beta=0.88, \lambda=2.25$



## COBRA: Anchoring Bias and Epsilon-Bounded Rationality

- ▶ “epsilon optimality”
- ▶ Anchoring bias: Full observation ( $\alpha = 0$ ) vs no observation ( $\alpha = 1$ )

$$\begin{aligned} & \max_{x,q,\gamma,a} \gamma \\ & s.t. x' = (1 - \alpha)x + \frac{\alpha}{N} \end{aligned}$$

$a$  is attacker's highest expected utility given  $x'$

$$q_j = 1 \text{ if } \text{AttEU}_j(x') \geq a - \epsilon$$

$$\gamma \leq \text{DefEU}_j(x) \text{ if } q_j = 1$$

- ▶ Experiments:  $\alpha = 0.37$  works best

## MATCH:Attacker aims to reduce the defender's utility

- ▶ Attacker may deviate from the best response to reduce the defender's expected utility
- ▶ Choose a target to maximize  
$$\frac{\text{Defender's utility loss due to deviation}}{\text{Adversary's utility loss due to deviation}}$$
- ▶ Defender: choose a strategy that maximize DefEU while bound the above value by  $\beta$
- ▶ Experiments:  $\beta = 1$

# Human Behavior Modeling & Learning

- ▶ Uncertainty and Bias Based Models
  - ▶ Prospect Theory [Kahneman and Tversky, 1979]
  - ▶ Anchoring bias and epsilon-bounded rationality [Pita et al, 2010]
  - ▶ Attacker aims to reduce the defender's utility [Pita et al, 2012]
- ▶ Quantal Response Based Models
  - ▶ Quantal Response [McKelvey and Palfrey, 1995]
  - ▶ Subjective Utility Quantal Response [Nguyen et al, 2013]
- ▶ Latest Models
  - ▶ Incorporating delayed observation [Fang et al, 2015]
  - ▶ Bounded rationality in repeated games [Kar et al, 2015]
  - ▶ Two-layered model [Nguyen et al, 2016]
  - ▶ Decision tree-based model [Kar & Ford et al, 2017]
- ▶ PAWS

# QR: Quantal Response Model

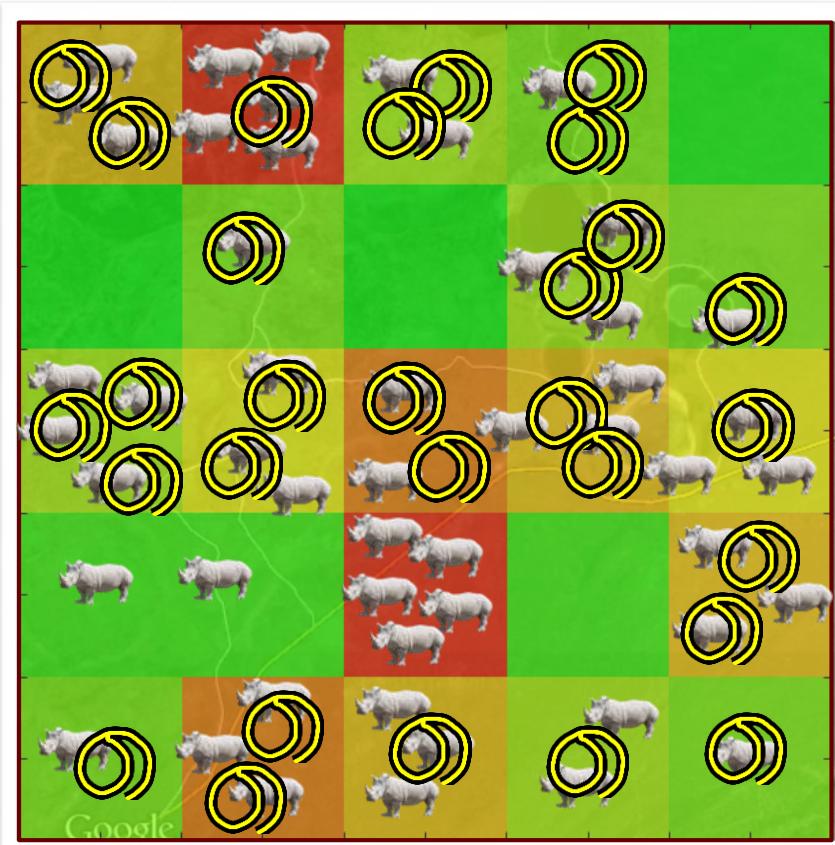
- ▶ Error in individual's response
  - ▶ Still: more likely to select better choices than worse choices
- ▶ Probability distribution of different responses
- ▶ Quantal best response:

$$q_j = \frac{e^{\lambda * \text{AttEU}_j(x)}}{\sum_i e^{\lambda * \text{AttEU}_i(x)}}$$

- ▶  $\lambda$ : represents error level ( $=0$  means uniform random)
  - ▶ Maximal likelihood estimation ( $\lambda=0.76$ )

# SUQR: Subjective Utility Quantal Response Model

►  $\text{SEU}_j = \sum_k w_k * f_j^k, q_j = \frac{e^{\lambda * \text{SEU}_j(x)}}{\sum_i e^{\lambda * \text{SEU}_i(x)}}$



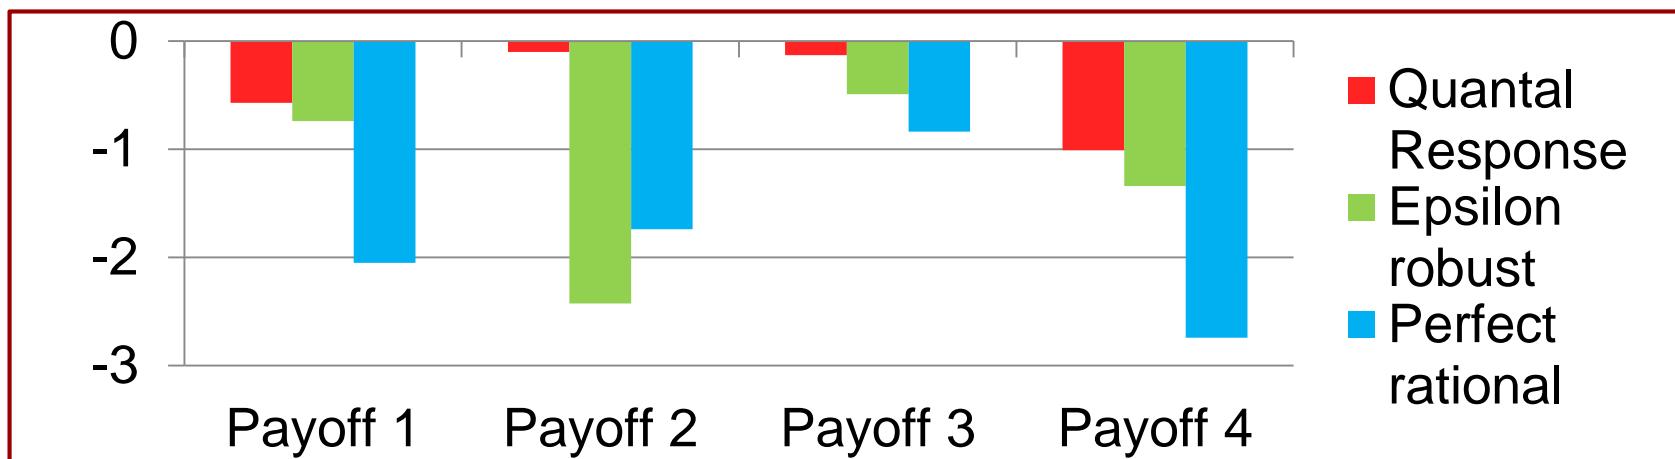
Coverage Probability  
+ Reward/Penalty

↓  
SUQR

Attack Probability

# Comparison of Model Performance

- ▶ Prospect Theory < DOBSS < COBRA < Quantal Response < MATCH < SUQR



MATCH wins	Draw	QR wins
42	52	6

MATCH wins	Draw	SUQR wins
1	8	13

# Human Behavior Modeling & Learning

- ▶ Uncertainty and Bias Based Models
  - ▶ Prospect Theory [Kahneman and Tversky, 1979]
  - ▶ Anchoring bias and epsilon-bounded rationality [Pita et al, 2010]
  - ▶ Attacker aims to reduce the defender's utility [Pita et al, 2012]
- ▶ Quantal Response Based Models
  - ▶ Quantal Response [McKelvey and Palfrey, 1995]
  - ▶ Subjective Utility Quantal Response [Nguyen et al, 2013]
- ▶ Latest Models
  - ▶ Incorporating delayed observation [Fang et al, 2015]
  - ▶ Bounded rationality in repeated games [Kar et al, 2015]
  - ▶ Two-layered model [Nguyen et al, 2016]
  - ▶ Decision tree-based model [Kar & Ford et al, 2017]
- ▶ PAWS Application

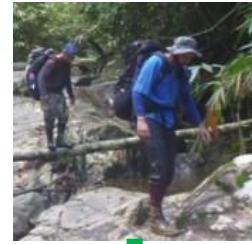
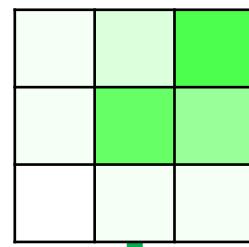
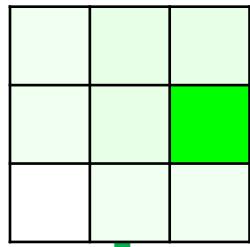
# GSG: Incorporating Delayed Observation



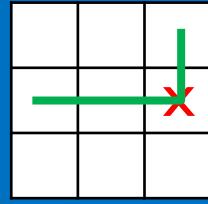
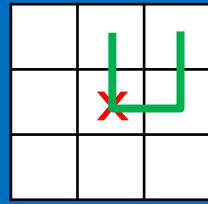
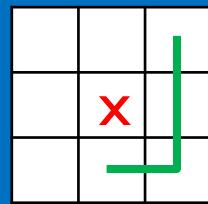
- ▶ Frequent and repeated attacks
  - ▶ Not one-shot / More data
- ▶ Attacker decision making
  - ▶ Limited surveillance / Less effort / Boundedly rational
- ▶ New model: Green Security Games

# GSG: Incorporating Delayed Observation

## Defender



Time

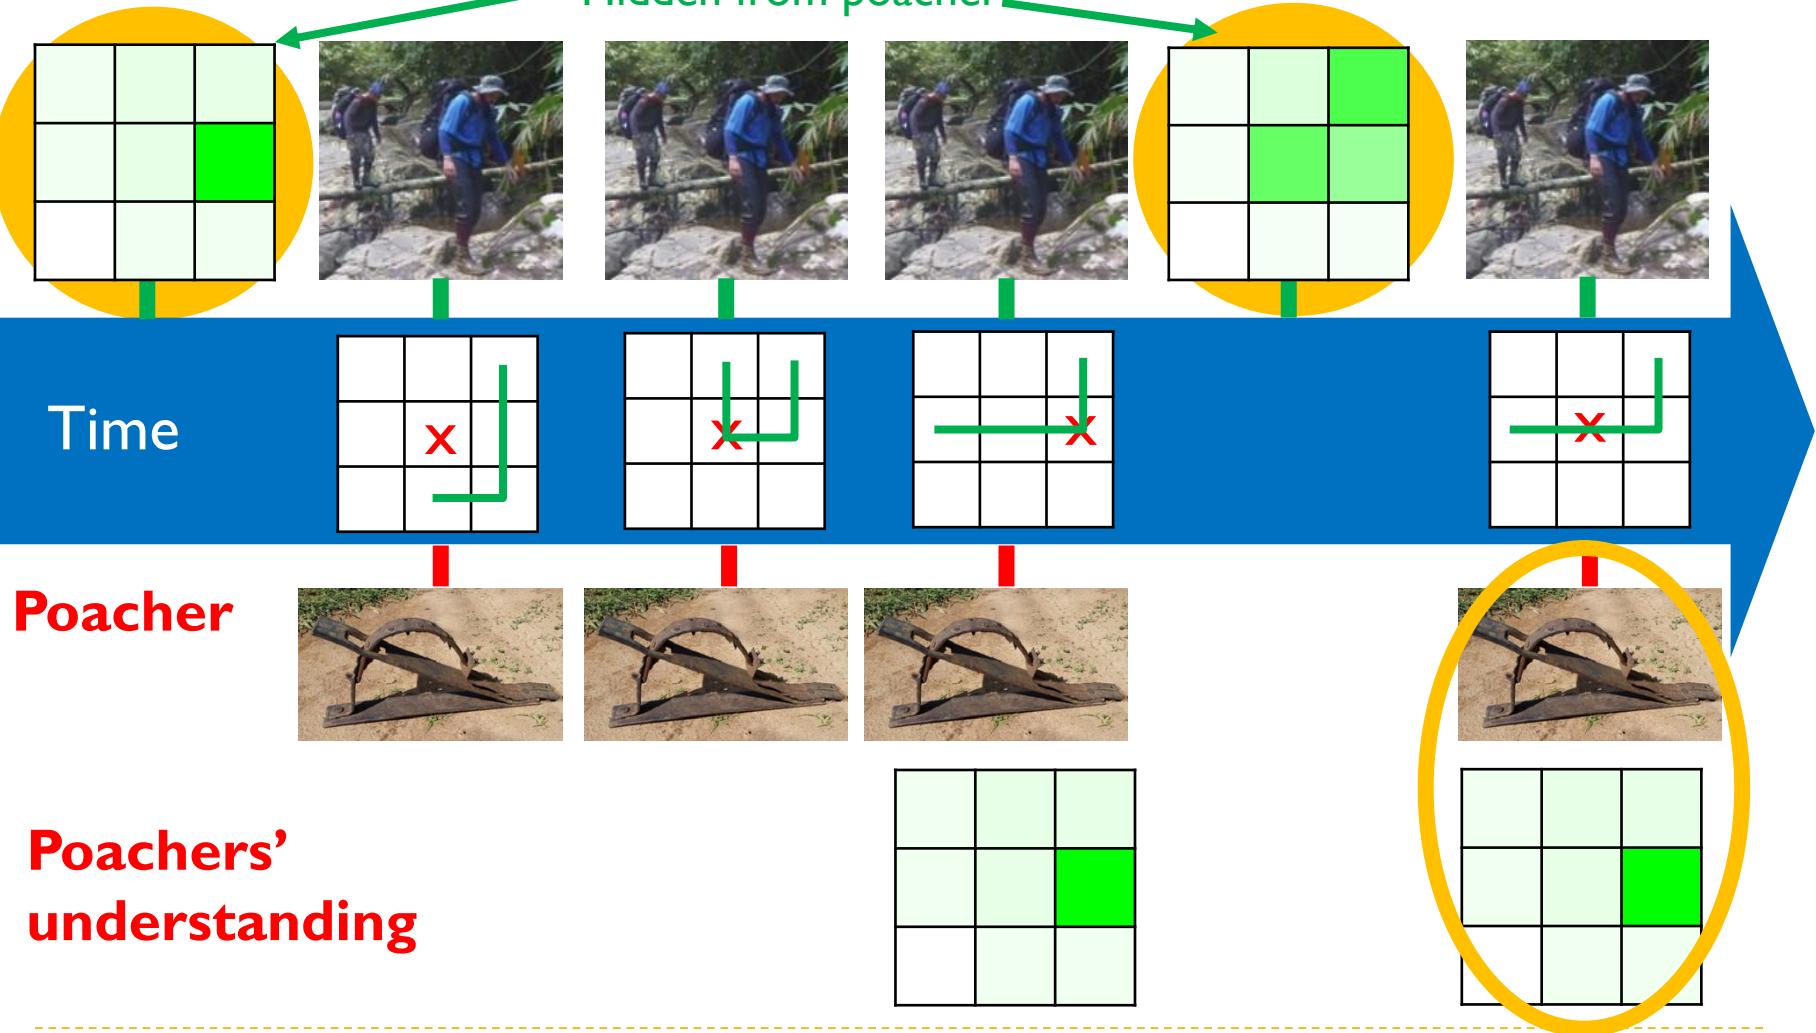


## Poacher



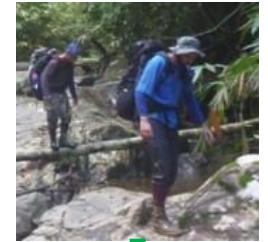
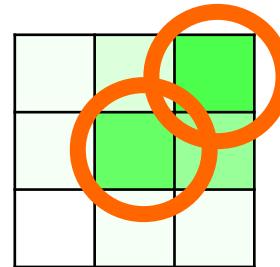
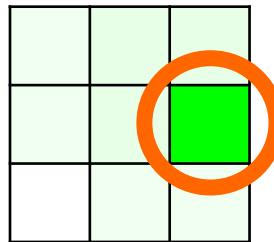
# GSG: Incorporating Delayed Observation

Defender

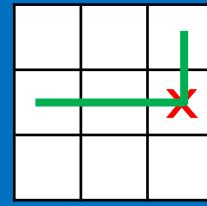
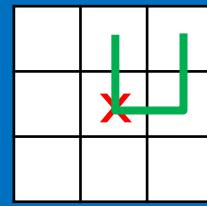
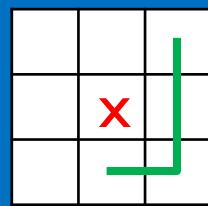


# GSG: Incorporating Delayed Observation

## Defender



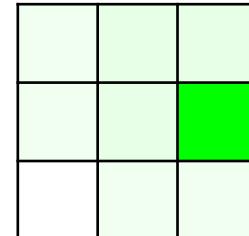
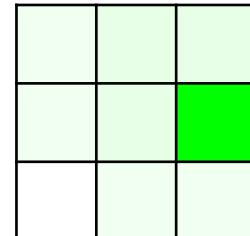
Time



Poacher



Poachers'  
understanding

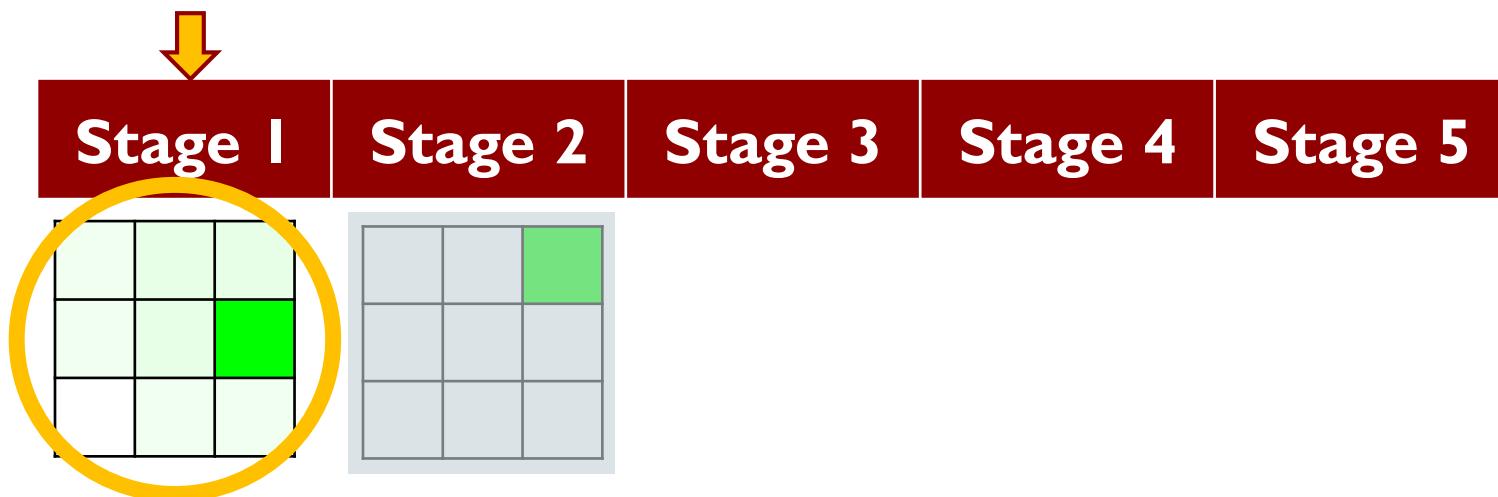


## GSG: Incorporating Delayed Observation

- ▶ A Green Security Game (GSG) is a  $T$  stage game where the defender protects  $N$  targets against  $L$  attackers. Defender chooses a mixed strategy  $c^t$  in stage  $t$ .
- ▶ A GSG attacker is characterized by his memory length  $\Gamma$ , coefficients  $\alpha_0, \dots, \alpha_\Gamma$  and SUQR model parameter  $\omega$ . In stage  $t$ , he responds to a convex combination of defender strategy in recent  $\Gamma + 1$  rounds:  $\eta_t = \sum_{\tau=0}^{\Gamma} \alpha_\tau c^{t-\tau}$

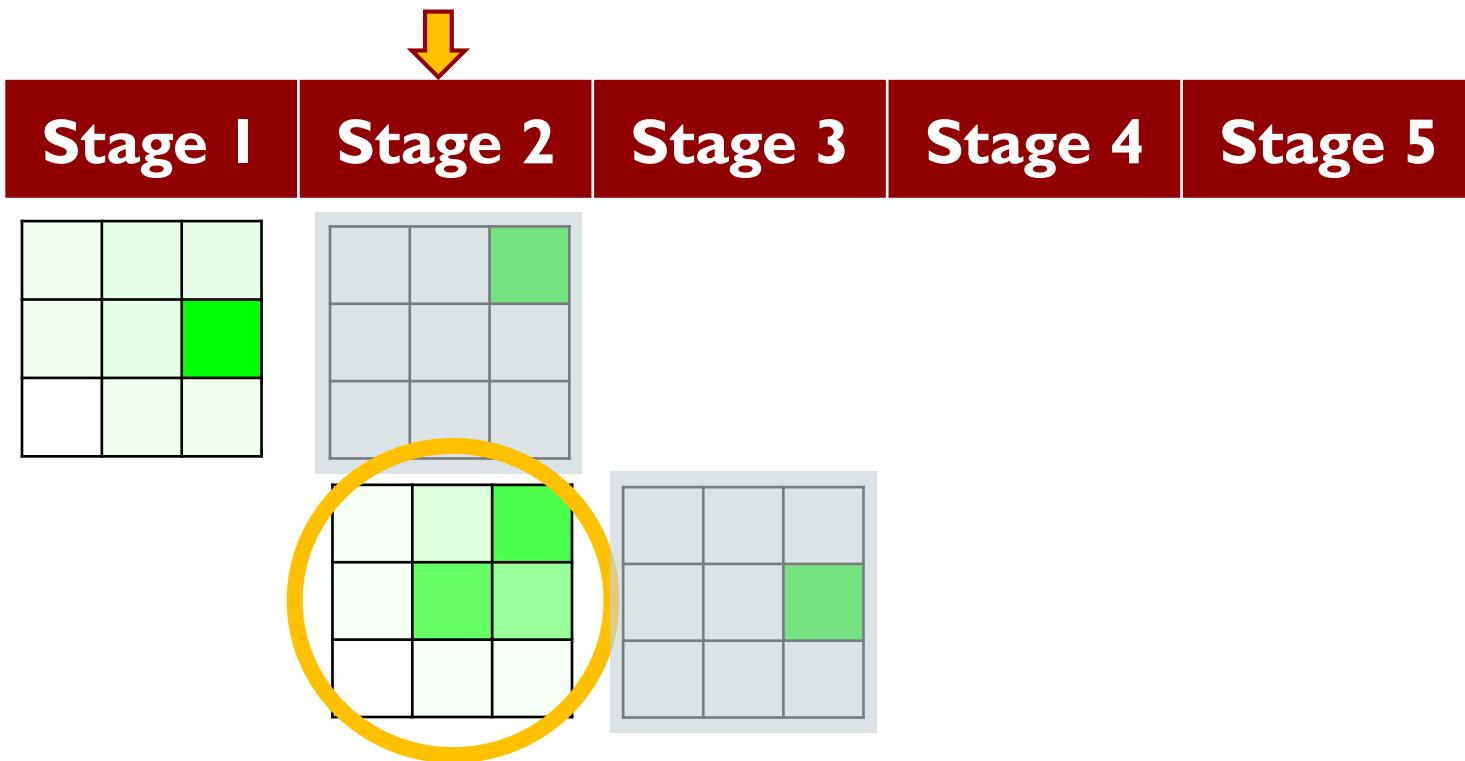
# GSG: Incorporating Delayed Observation

- ▶ Plan Ahead – M (PA-M)
- ▶ Plan ahead M stages



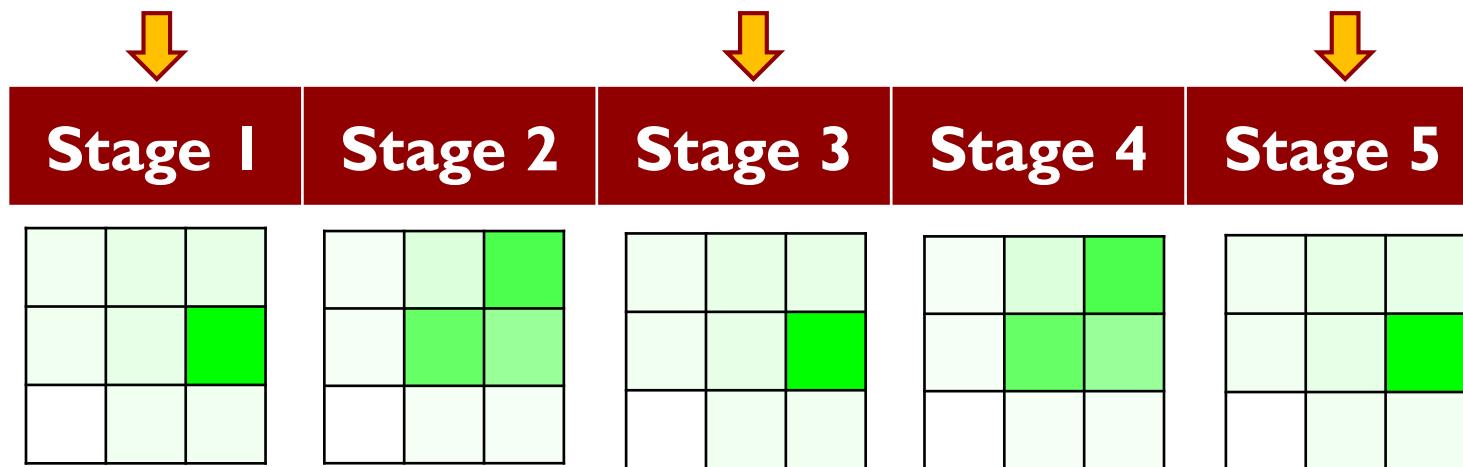
# GSG: Incorporating Delayed Observation

- ▶ Plan Ahead – M (PA-M)
- ▶ Plan ahead M stages

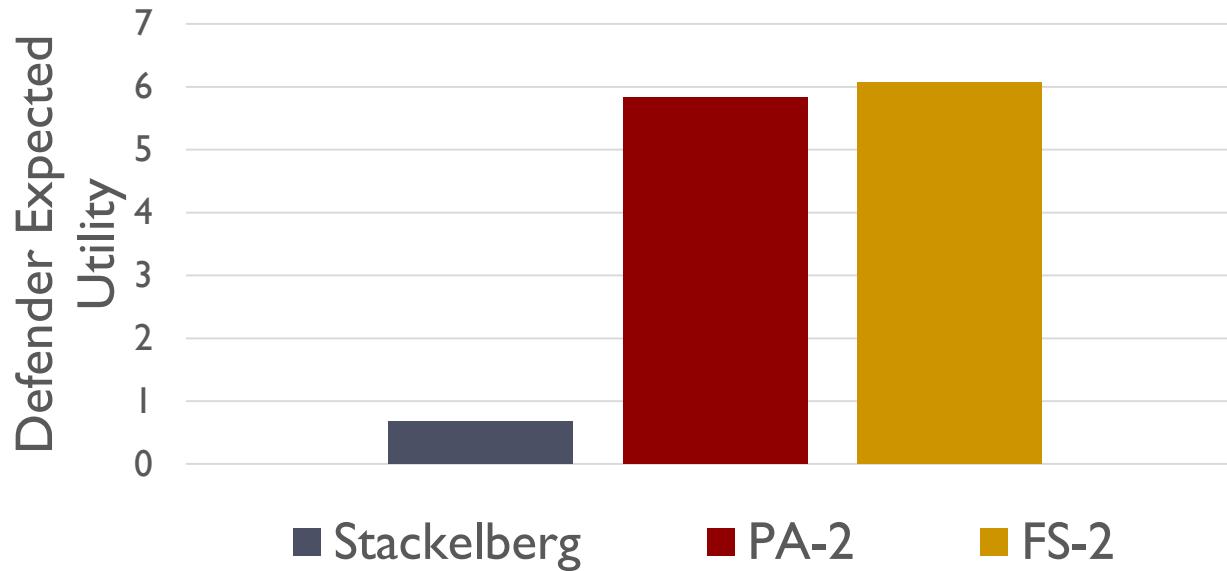


# GSG: Incorporating Delayed Observation

- ▶ An alternative: Fixed Sequence – M (FS-M)
- ▶ Use M strategies repeatedly



# GSG: Incorporating Delayed Observation



- ▶ **Theorem 3:** In a GSG with  $T$  rounds, for  $\Gamma < M \leq T$ , there exists a cyclic defender strategy profile  $[s]$  with period  $M$  that is a  $(1 - \frac{\Gamma}{T}) \frac{Z-1}{Z+1}$  approximation of the optimal strategy profile in terms of the normalized utility, where  $Z = \lceil \frac{T-\Gamma+1}{M} \rceil$

# Human Behavior Modeling & Learning

- ▶ Uncertainty and Bias Based Models
  - ▶ Prospect Theory [Kahneman and Tversky, 1979]
  - ▶ Anchoring bias and epsilon-bounded rationality [Pita et al, 2010]
  - ▶ Attacker aims to reduce the defender's utility [Pita et al, 2012]
- ▶ Quantal Response Based Models
  - ▶ Quantal Response [McKelvey and Palfrey, 1995]
  - ▶ Subjective Utility Quantal Response [Nguyen et al, 2013]
- ▶ Latest Models
  - ▶ Incorporating delayed observation [Fang et al, 2015]
  - ▶ Bounded rationality in repeated games [Kar et al, 2015]
  - ▶ Two-layered model [Nguyen et al, 2016]
  - ▶ Decision tree-based model [Kar & Ford et al, 2017]
- ▶ PAWS Application

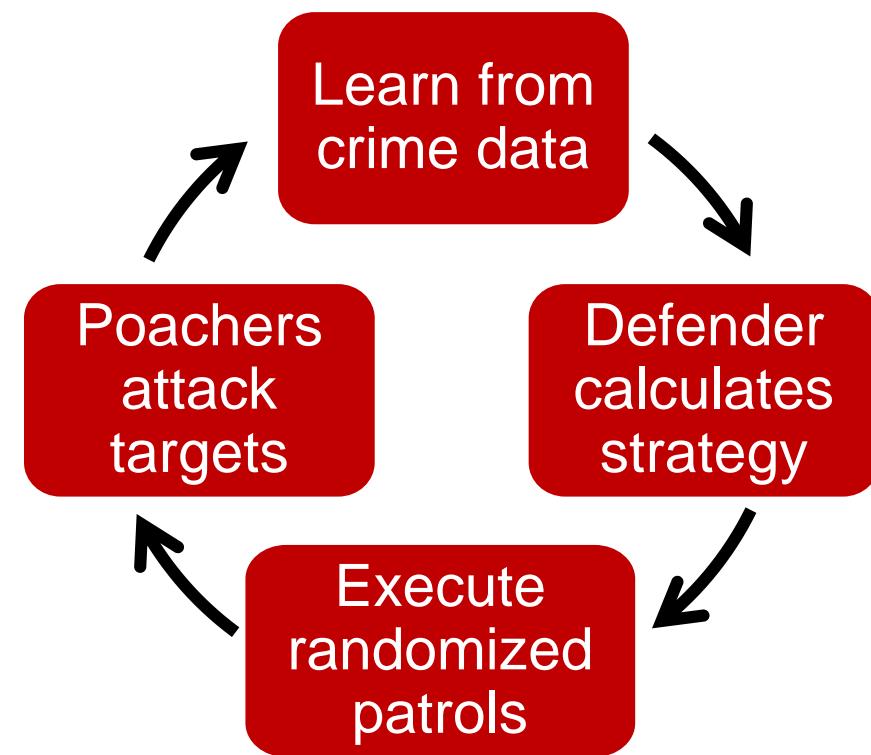
# SHARP: Bounded Rationality in Repeated Games



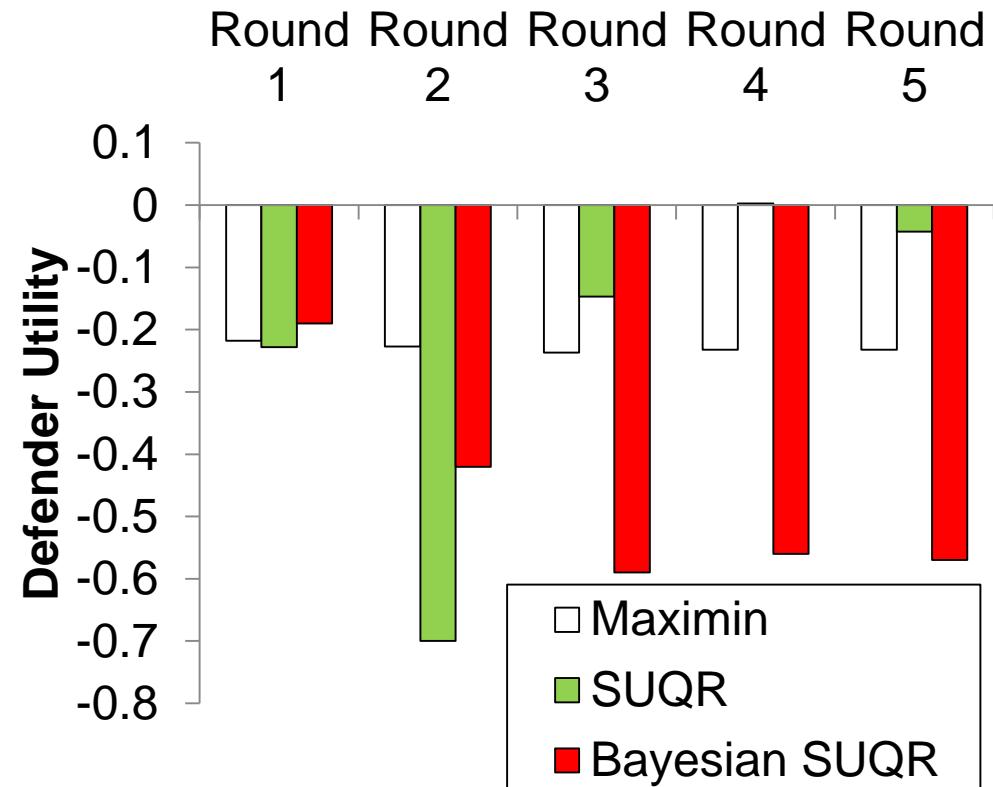
Game 4

Total: \$1.5

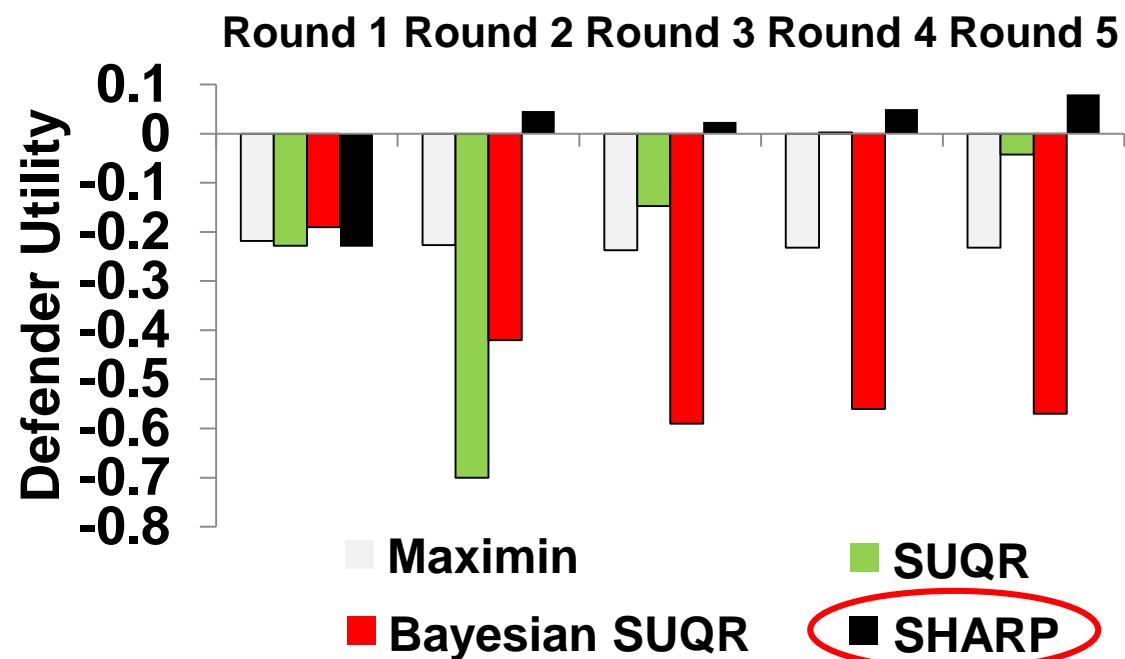
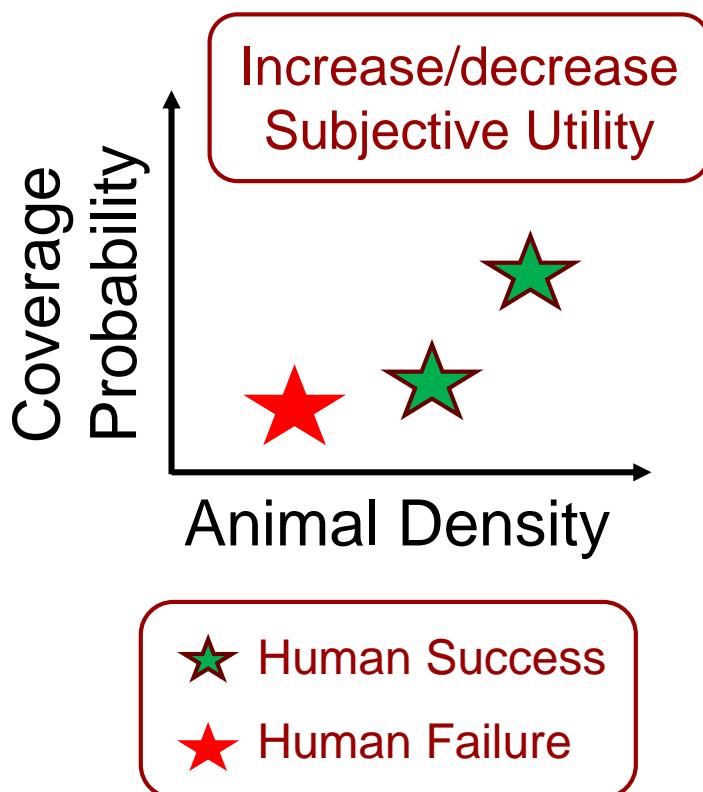
# SHARP: Bounded Rationality in Repeated Games



Repeated games on AMT: 35 weeks, 40 human subjects 10,000 emails!

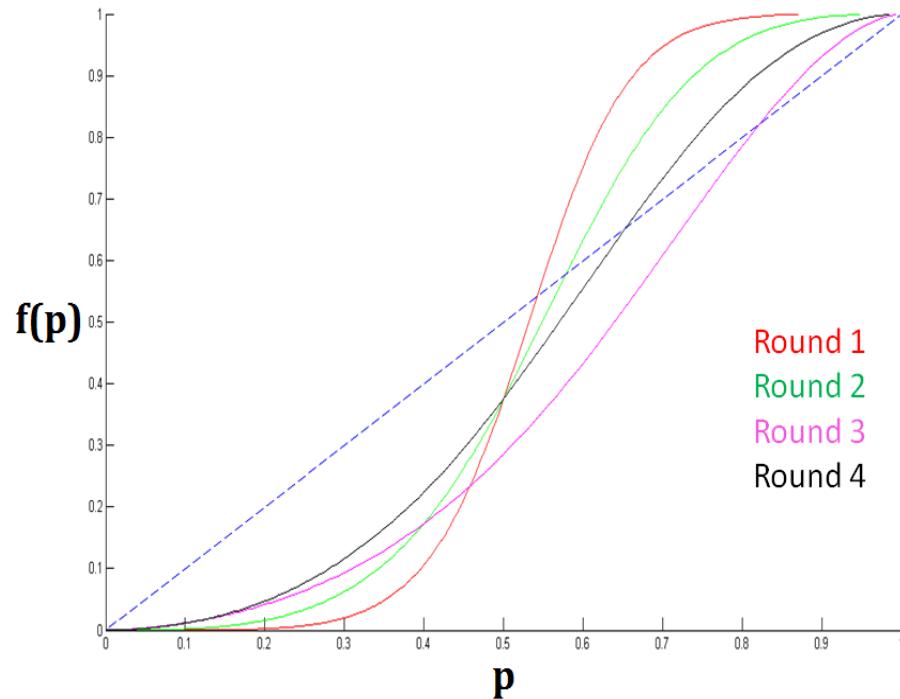


# SHARP: Bounded Rationality in Repeated Games

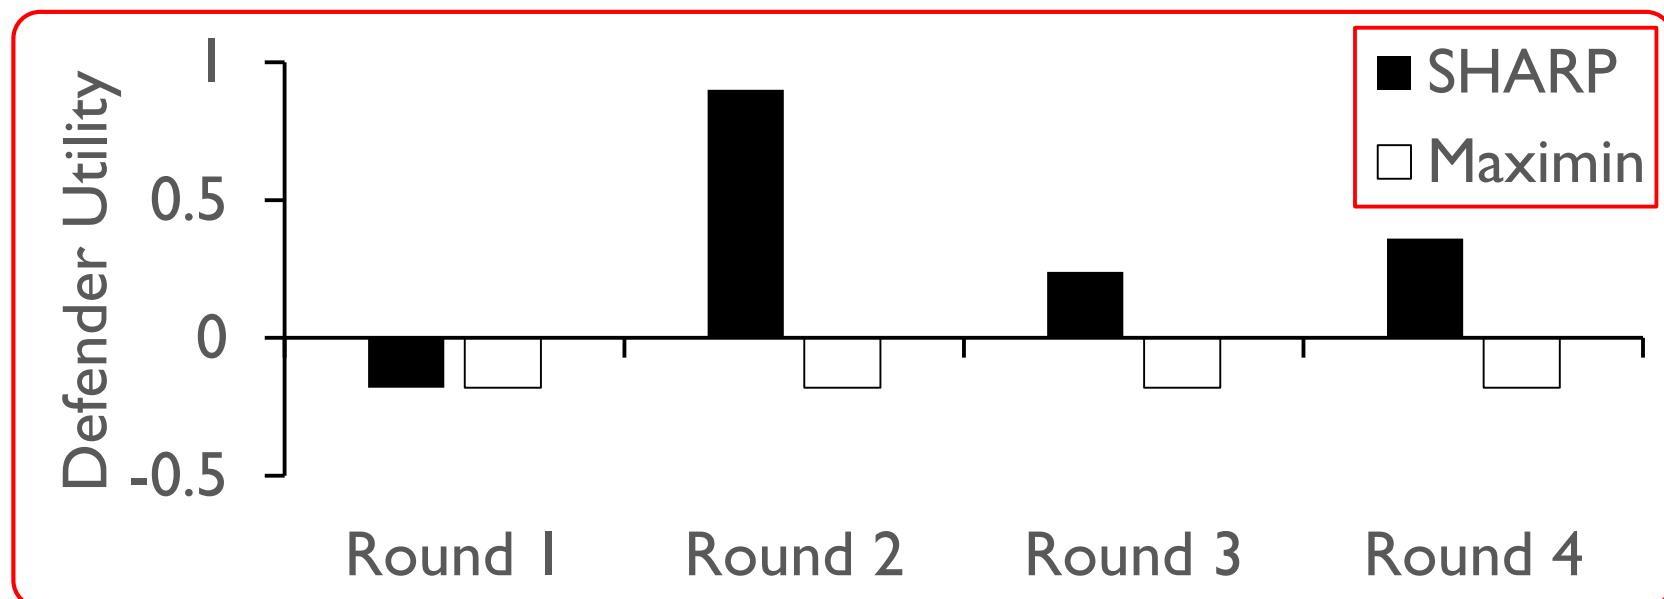


# SHARP: Bounded Rationality in Repeated Games

- ▶ Adversary's probability weighting function is S-shaped.
  - ▶ Contrary to Prospect Theory



# SHARP: Bounded Rationality in Repeated Games



# Human Behavior Modeling & Learning

- ▶ Uncertainty and Bias Based Models
  - ▶ Prospect Theory [Kahneman and Tversky, 1979]
  - ▶ Anchoring bias and epsilon-bounded rationality [Pita et al, 2010]
  - ▶ Attacker aims to reduce the defender's utility [Pita et al, 2012]
- ▶ Quantal Response Based Models
  - ▶ Quantal Response [McKelvey and Palfrey, 1995]
  - ▶ Subjective Utility Quantal Response [Nguyen et al, 2013]
- ▶ Latest Models
  - ▶ Incorporating delayed observation [Fang et al, 2015]
  - ▶ Bounded rationality in repeated games [Kar et al, 2015]
  - ▶ Two-layered model [Nguyen et al, 2016]
  - ▶ Decision tree-based model [Kar & Ford et al, 2017]
- ▶ PAWS Application

# Real-World Data

- ▶ Queen Elizabeth National Park

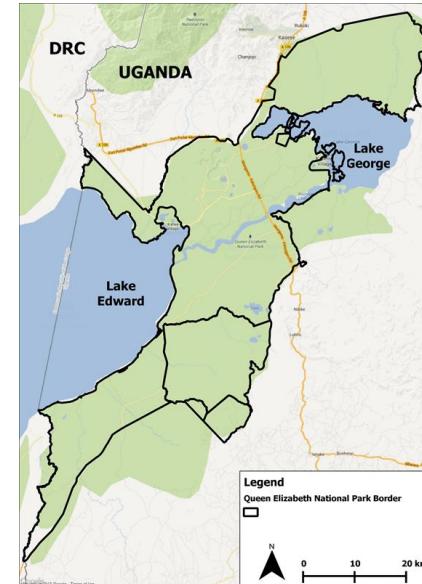
- ▶ 1,978 sq. km
- ▶ 2003-2015

- ▶ Geospatial Features

- ▶ Terrain (e.g., forest, slope)
- ▶ Distance to {Town, Water, Outpost}

- ▶ Ranger Coverage

- ▶ Crime Observations



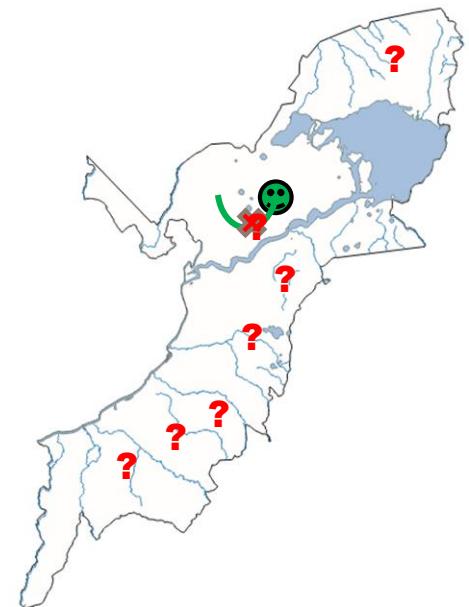
# Real-World Data: Challenges

## ▶ “Missing” poaching data

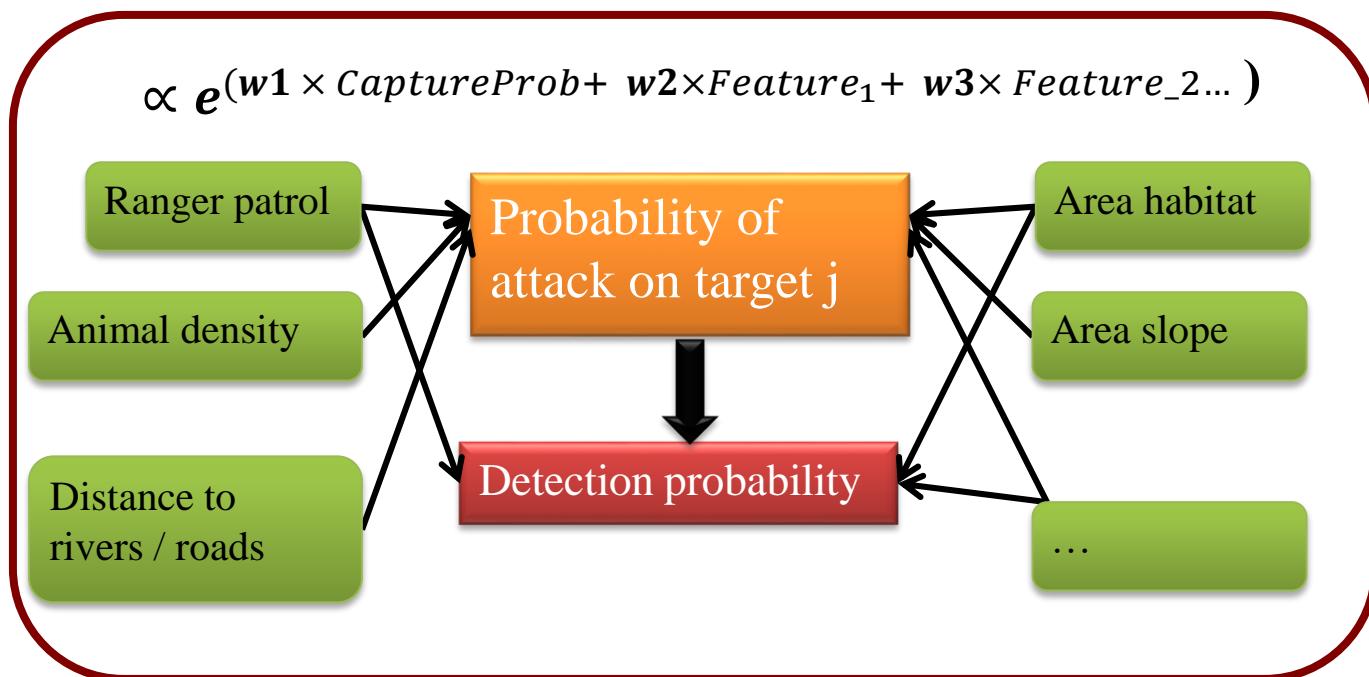
- ▶ Limited patrol resources  
(silent victims)
- ▶ Imperfect observations  
(e.g., hidden snares)

## ▶ Consequences

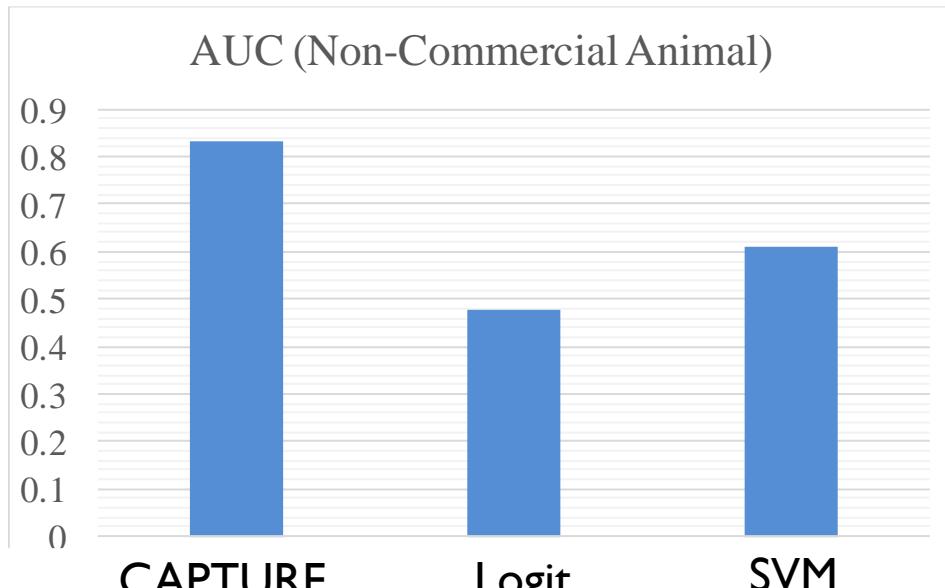
- ▶ Uncertainty in negative labels
- ▶ Class imbalance



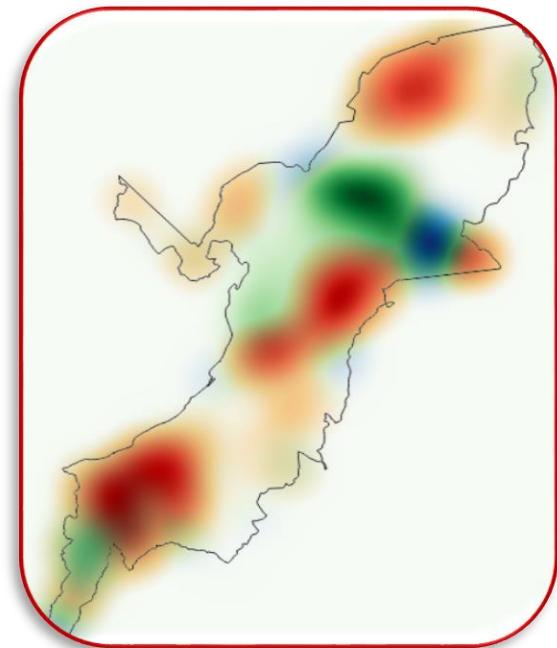
# CAPTURE:Two-Layered Model



# CAPTURE: Two-Layered Model



Dry Season (June-August 2008)



Green – Animal Density;  
Blue – Defender Strategy;  
Red – Observed Attack Probability

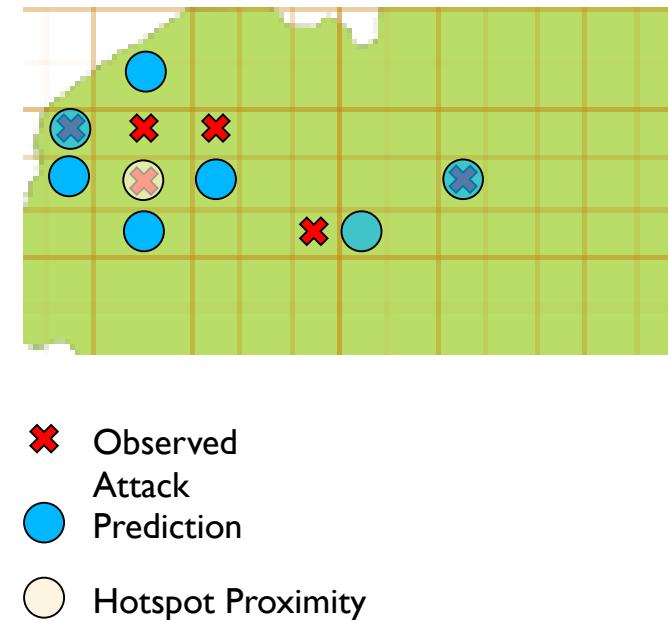
# CAPTURE: Two-Layered Model

- ▶ Limitations
  - ▶ Limited Data
  - ▶ Predicting Everywhere
  - ▶ Slow Learning
- ▶ Variations
  - ▶ Simpler observation layer
  - ▶ Use previous coverage instead of current coverage
  - ▶ Exponentially penalize the attractiveness of inaccessible areas
  - ▶ No improvement

# INTERCEPT: Decision-Tree Based Model

- ▶ BoostIT
- ▶ Consider spatial correlations (hotspots)

- ▶ Learn a decision tree
- ▶ Compute predictions
- ▶ Hotspot proximity computation
- ▶ Learn a new decision tree with hotspot proximity as a feature
- ▶ Repeat until a stopping condition is reached

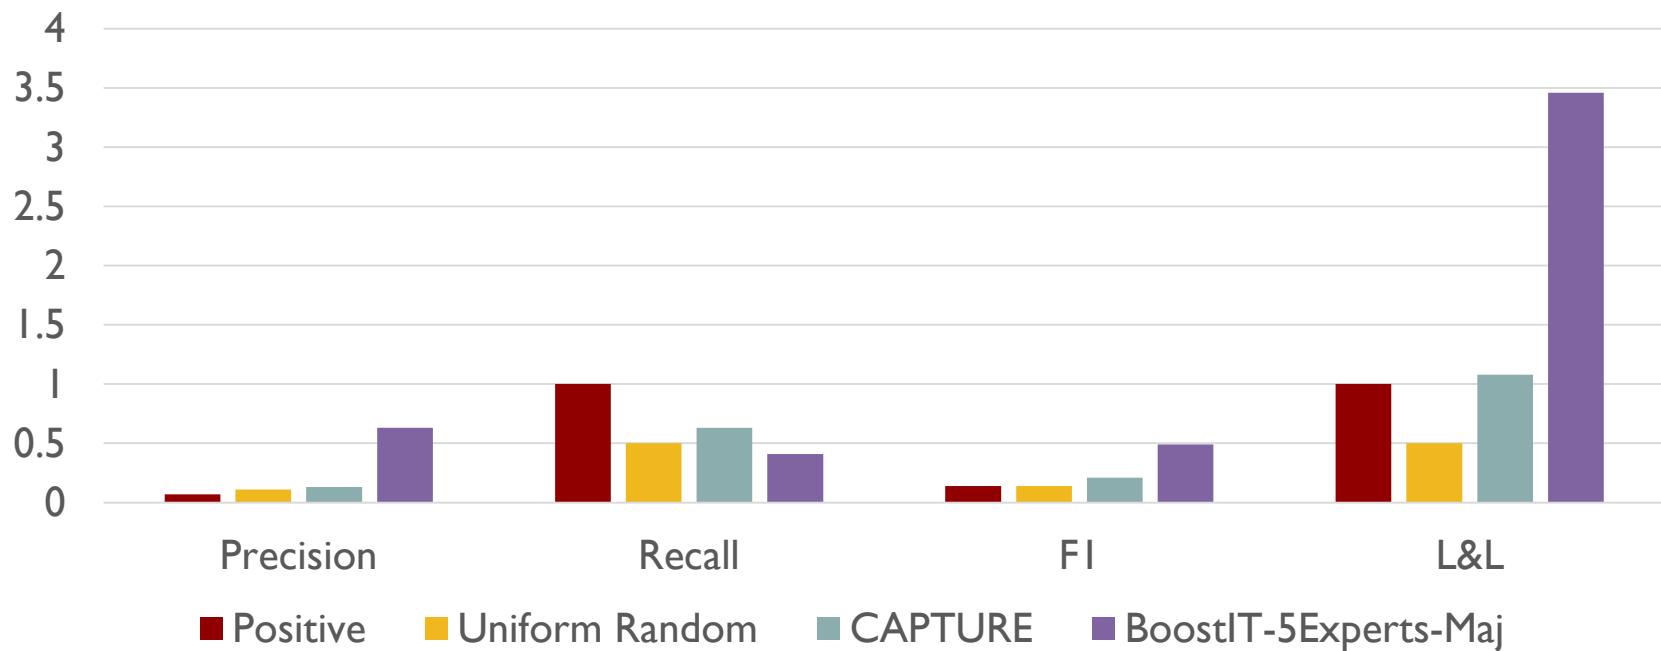


# INTERCEPT: Decision-Tree Based Model

- ▶ Extensive Empirical Evaluation
  - ▶ 40 models w/ total of **192 model variations**
  - ▶ Decision tree ensemble outperforms all other (complex) models including CAPTURE
    - ▶ Standard decision tree, 2 BoostITs ( $\alpha = 2, 3$ ), 2 Decision Trees (FP cost = 0.6, 0.9)
  - ▶ Fast runtime and interpretability led to its deployment (obstacles for CAPTURE)
- ▶ Datasets
  - ▶ Trained: 2003-2014, Tested: 2015
  - ▶ Trained: 2003-2013, Tested: 2014

# INTERCEPT: Decision-Tree Based Model

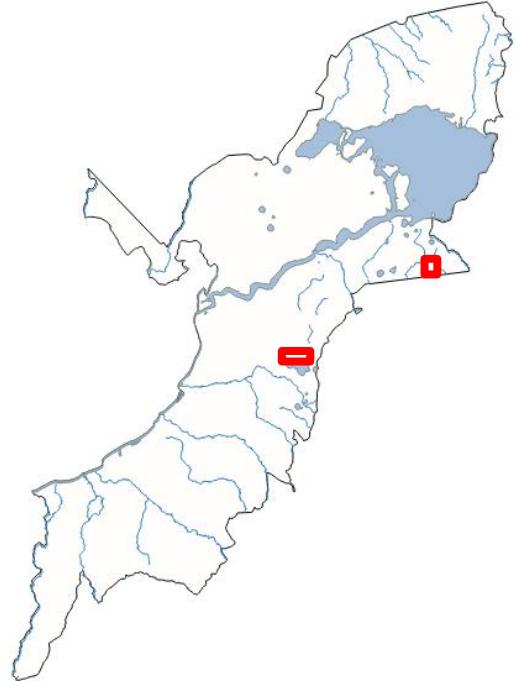
## Empirical Evaluation: Key Attack Prediction Results: 2015



$$L\&L(D, T_e) = \frac{recall^2}{\Pr[f(T_e) = 1]}$$

# Real-world Deployment: Field Test I (1 month)

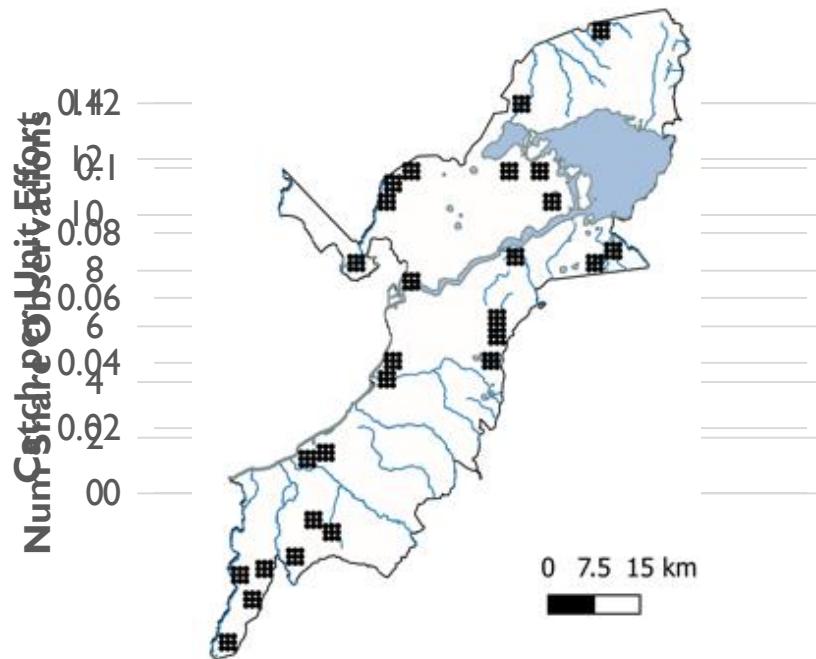
- ▶ Trespassing
  - ▶ 19 signs of litter, ashes, etc.
- ▶ Poached animals
  - ▶ 1 poached elephant
- ▶ Snaring
  - ▶ 1 active snare
  - ▶ 1 cache of 10 antelope snares
  - ▶ 1 roll of elephant snares
- ▶ Snaring hit rates
  - ▶ Outperform 91% of months



Historical Base Hit Rate	Our Hit Rate
Average: 0.73	3

# Real-world Deployment: Field Test 2 (5 months)

- ▶ 27 areas (9-sq km each)
- ▶ 454 km patrolled in total
- ▶ 2 experiment groups
  - ▶ 1:  $\geq 50\%$  attack prediction rate
    - ▶ 5 areas
  - ▶ 2:  $< 50\%$  attack prediction rate
    - ▶ 22 areas
- ▶ Catch Per Unit Effort (CPUE)
  - ▶ Unit Effort = km walked



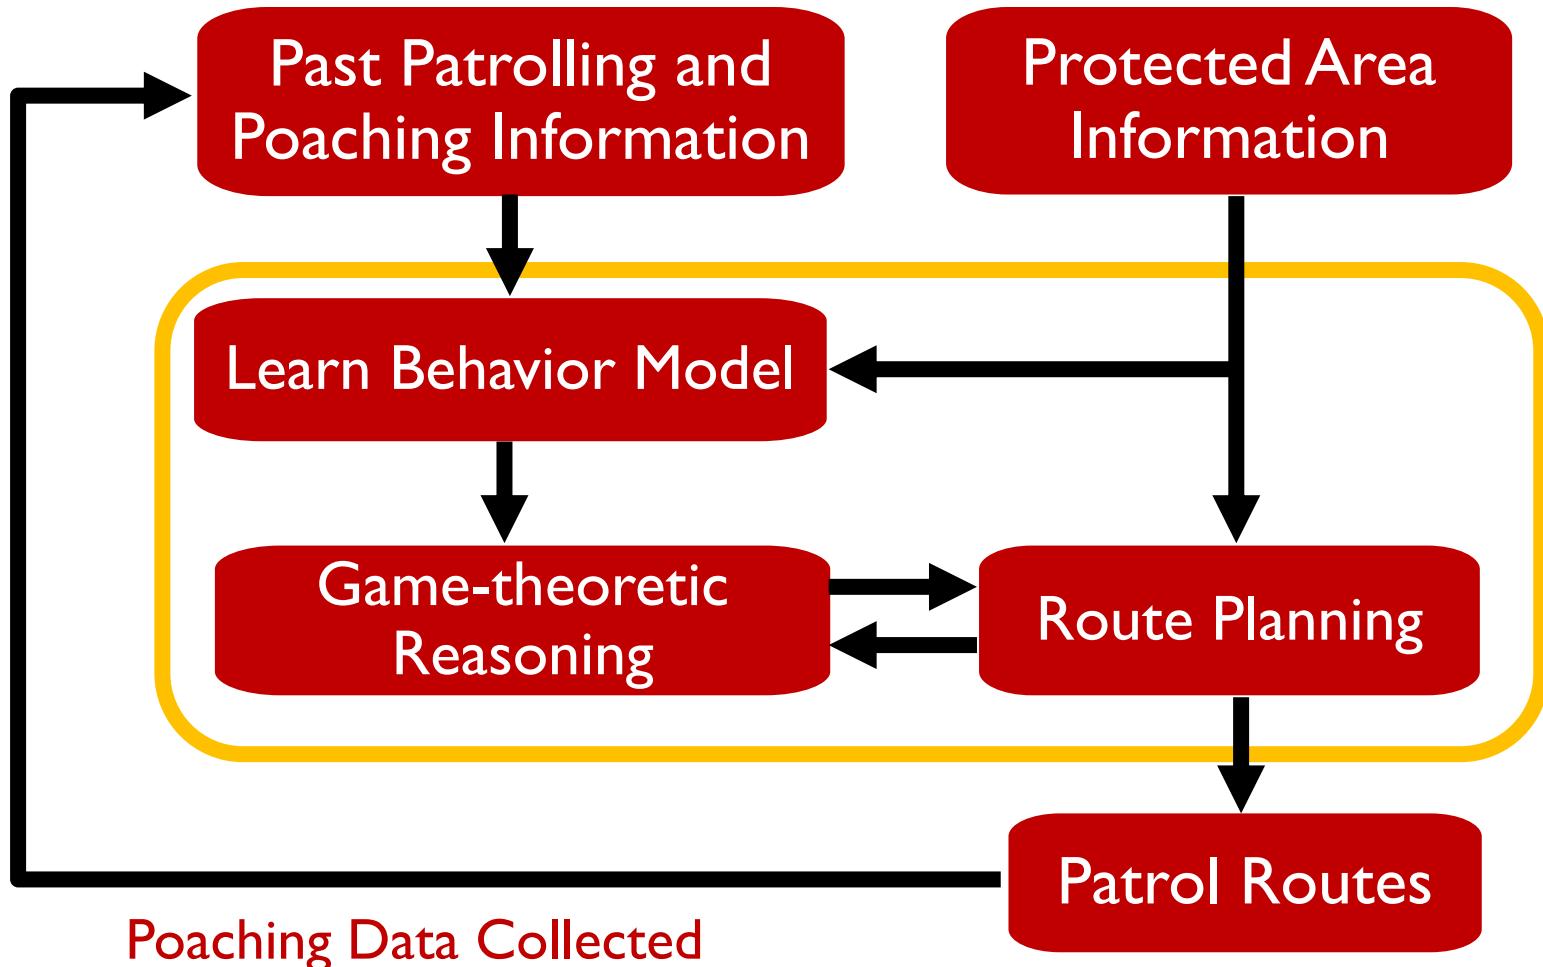
# Deployment Results Summary

- ▶ First Field Test: Demonstrated potential for predictive analytics in the field
- ▶ Second Field Test: Demonstrated the selectiveness of our model in the field
- ▶ We saved animals!

# Human Behavior Modeling & Learning

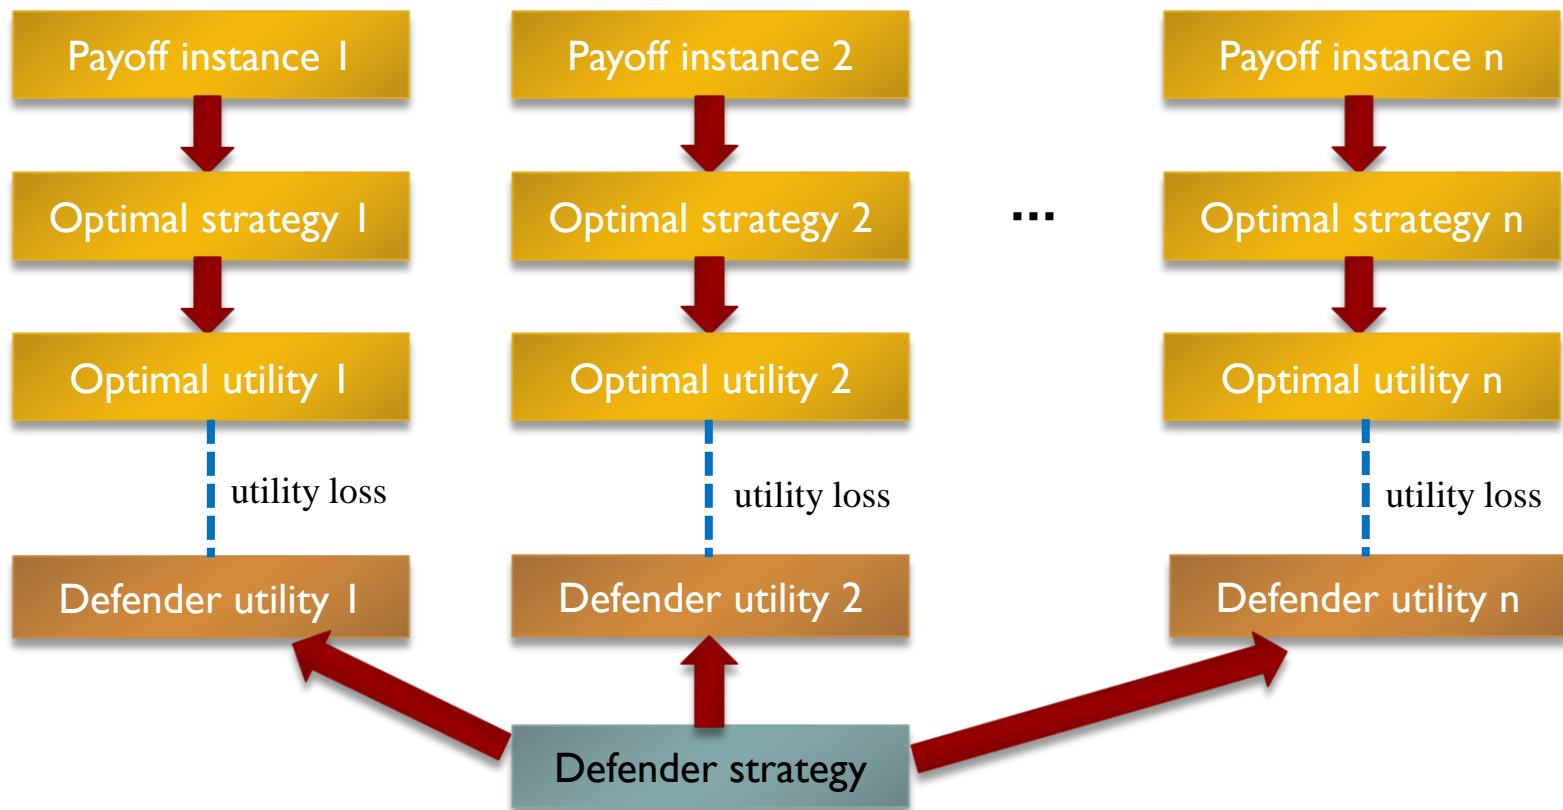
- ▶ Uncertainty and Bias Based Models
  - ▶ Prospect Theory [Kahneman and Tversky, 1979]
  - ▶ Anchoring bias and epsilon-bounded rationality [Pita et al, 2010]
  - ▶ Attacker aims to reduce the defender's utility [Pita et al, 2012]
- ▶ Quantal Response Based Models
  - ▶ Quantal Response [McKelvey and Palfrey, 1995]
  - ▶ Subjective Utility Quantal Response [Nguyen et al, 2013]
- ▶ Latest Models
  - ▶ Incorporating delayed observation [Fang et al, 2015]
  - ▶ Bounded rationality in repeated games [Kar et al, 2015]
  - ▶ Two-layered model [Nguyen et al, 2016]
  - ▶ Decision tree-based model [Kar & Ford et al, 2017]
- ▶ PAWS

# PAWS Overview



# PAWS Component: Game-theoretic Reasoning

- ▶ Handling payoff uncertainty
- ▶ Behavioral minimax regret [Nguyen et al, 2016]



# PAWS Component: Game-theoretic Reasoning

## ► Uncertainty interval

		Adversary	
		Target 1	Target 2
Defender	Target 1	4, [-4, -2]	-1, [0, 2]
	Target 2	-5, [4, 6]	2, [-2, 0]

# PAWS Component: Game-theoretic Reasoning

- ▶ Defender's utility of playing "x" given payoff instance
  - ▶ Sampled one payoff instance from interval

■  $\text{defUtility}(x) = -0.5$

		Adversary		
		Target 1	Target 2	x
		Target 1	4, -4	-1, 1
		Target 2	-5, 5	2, -2
Defender	q		0.6	0.4

# PAWS Component: Game-theoretic Reasoning

- ▶ Defender's regret for playing "x" given payoff instance
  - ▶ Utility loss of defender to play x compared to optimal.

- $\text{defUtility}(x) = -0.5$
- $\text{defUtility}(x^*) = 0.2$
- $\text{regret}(x) = 0.7$

Defender

		Adversary			
		Target 1	Target 2	x	$x^{*,l}$
Target 1		4, -4	-1, 1	0.3	0.7
Target 2		-5, 5	2, -2	0.7	0.3
q		0.6	0.4		
q'		0.4	0.6		

# PAWS Component: Game-theoretic Reasoning

- ▶ Defender regret of “ $x$ ” higher for another payoff instance
  - ▶ Max regret: Max over all payoff instances under uncertainty

- $\text{defUtility}(x) = -0.9$
- $\text{defUtility}(x^*) = 0.3$
- $\text{regret}(x) = 1.2$

**max\_regret( $x$ ) = 1.2**

Defender

		Adversary			
		Target 1	Target 2	$x$	$x^{*,2}$
Target 1		6, -6	-1, 1	0.3	0.8
Target 2		-4, 4	0, 0	0.7	0.2
$q$		0.8	0.2		
$q^2$		0.2	0.8		

# PAWS Component: Game-theoretic Reasoning

- ▶ Objective
  - ▶ Compute optimal defender strategy that minimizes max regret.
- ▶ Optimization

Infinite #regret constraints

$$\begin{aligned} & \min_{x \in X, r \in R} r \\ \text{s.t. } & r \geq \text{regret}(x, \text{payoff}), \quad \text{" payoff } \hat{\in} I \end{aligned}$$

A red arrow points from the text "Utility loss of playing x given attacker follows SUQR" up towards the constraint "r ≥ regret(x, payoff)".

Utility loss of playing  $x$   
given attacker follows SUQR

# PAWS Component: Game-theoretic Reasoning

## ► ARROW: Incremental Payoff Generation

$$\min_{x \in X, r \in R} r$$

s.t.  $r \geq \text{regret}(x, \text{payoff})$ , " payoff  $\in I$

### Master:

- **Idea:** solve relaxed MMR with small sample set  $S$  of attacker payoffs, obtain **LB**

### Slave:

- **Idea:** find new attacker payoff to add to the sample set, obtain **UB**

# PAWS Component: Game-theoretic Reasoning

x <sup>1</sup>	
Target 1	0.3
Target 2	0.7

		Payoff 1	
		Target 1	Target 2
Target 1	Target 1	4, -4	-1, 1
	Target 2	-5, 5	2, -2

# PAWS Component: Game-theoretic Reasoning

	$x^1$	$x^{1,2}$
Target 1	0.3	0.4
Target 2	0.7	0.6

		Payoff 1	
		Target 1	Target 2
Target 1	Target 1	4, -4	-1, 1
	Target 2	-5, 5	2, -2

		Payoff 2	
		Target 1	Target 2
Target 1	Target 1	3, -3	0, 0
	Target 2	-4, 4	1, -1

# PAWS Component: Game-theoretic Reasoning

	$x^1$	$x^{1,2}$	$x^{1,2,3}$	...
Target 1	0.3	0.4	0.6	...
Target 2	0.7	0.6	0.4	...

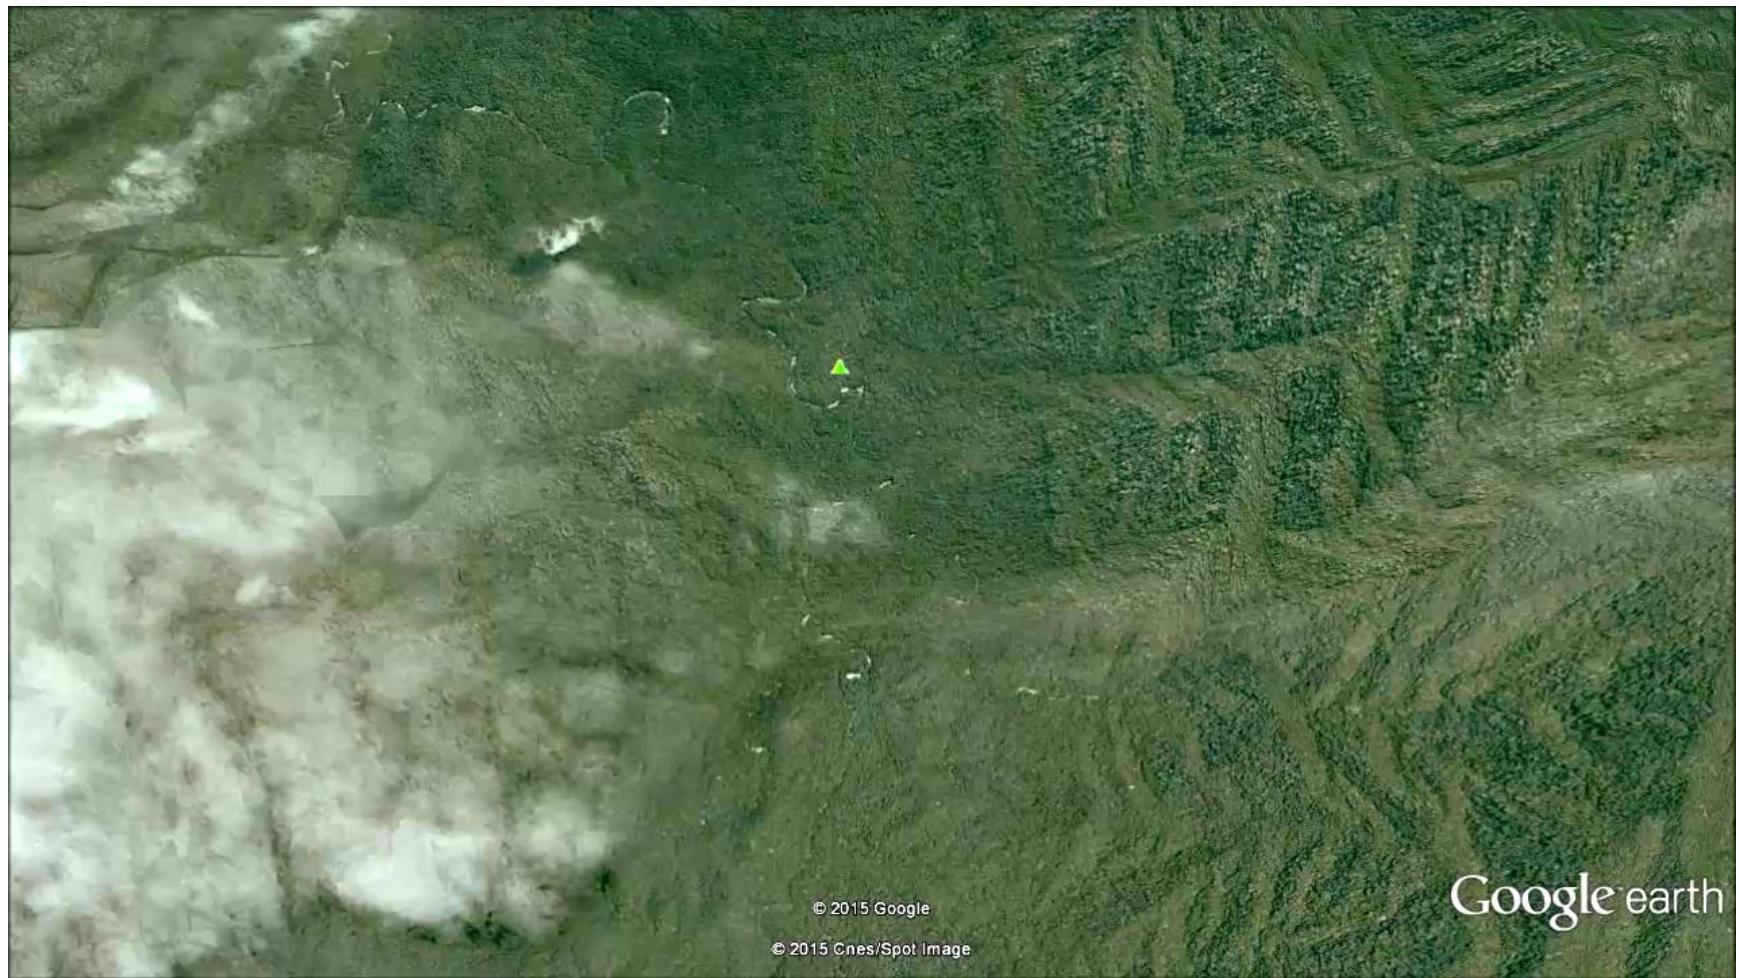
	Payoff 1	
	Target 1	Target 2
Target 1	4, -4	-1, 1
Target 2	-5, 5	2, -2

	Payoff 2	
	Target 1	Target 2
Target 1	3, -3	0, 0
Target 2	-4, 4	1, -1

	Payoff 3	
	Target 1	Target 2
Target 1	5, -5	2, -2
Target 2	6, -6	0, 0

...

# PAWS Component: Route Planning



Google earth

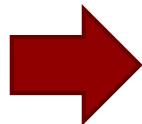
# PAWS Component: Route Planning

Defender Strategy

0.03	0.05	0.15
0.03	0.28	0.40
0.00	0.03	0.03



Patrol Route (2D)

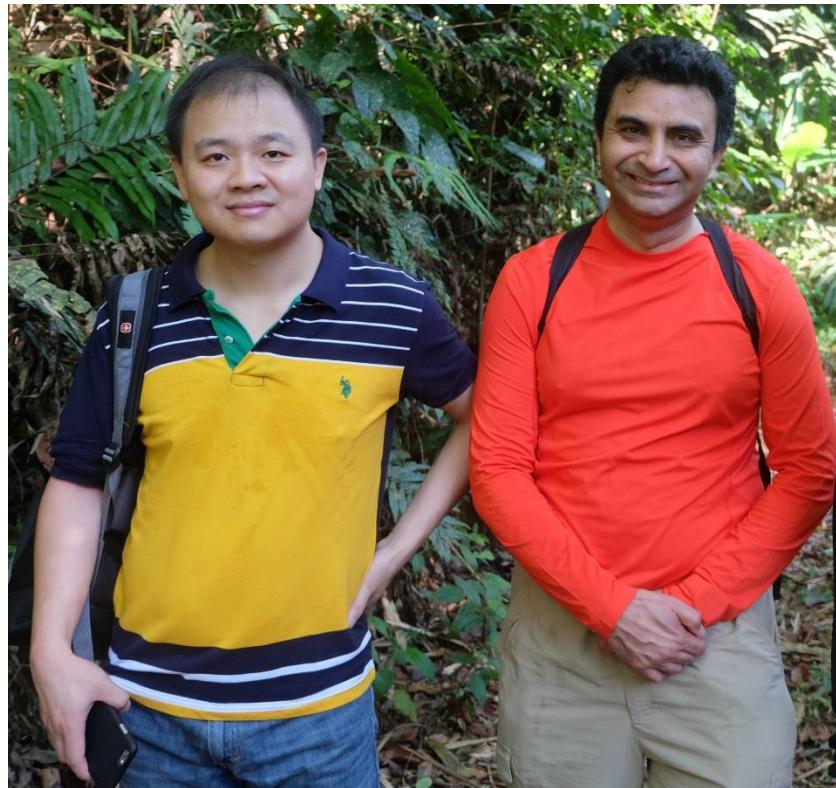


Patrol Route (3D)



# PAWS Component: Route Planning

- ▶ 8-hour patrol in April 2015: patrolling is not easy!

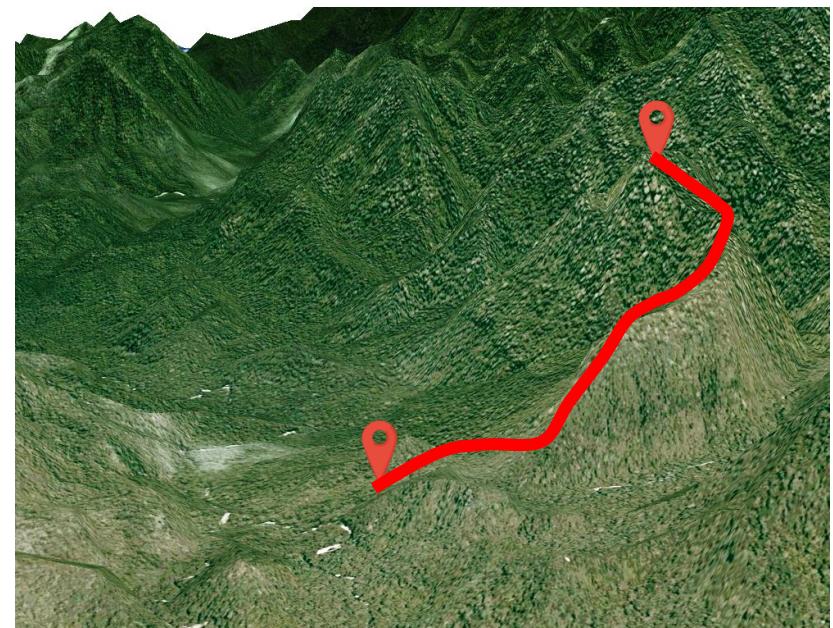
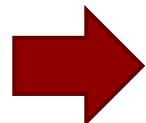
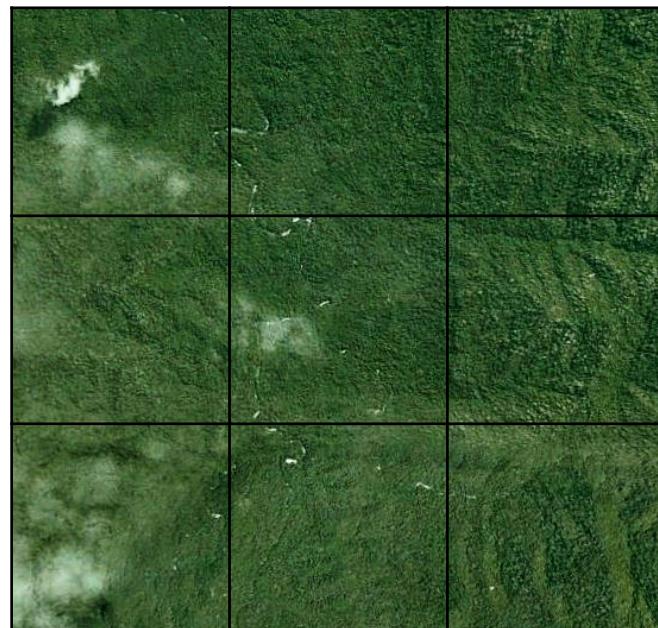


# PAWS Component: Route Planning



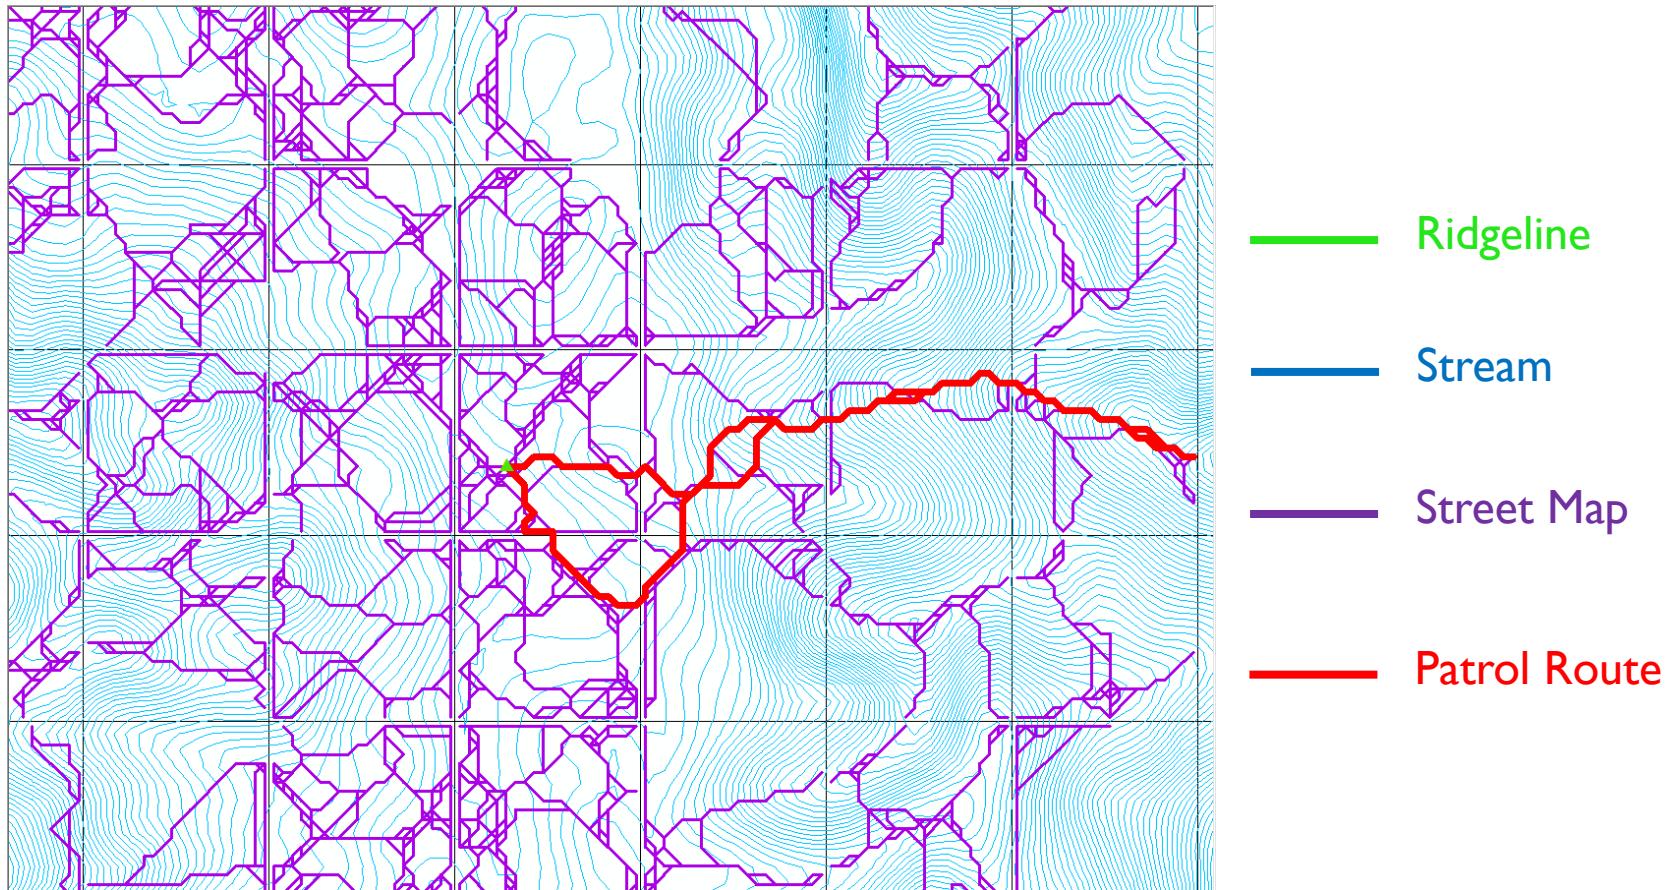
# PAWS Component: Route Planning

- ▶ Grid based → Route based
- ▶ Hierarchical modeling: Focus on terrain features
- ▶ Build virtual street map

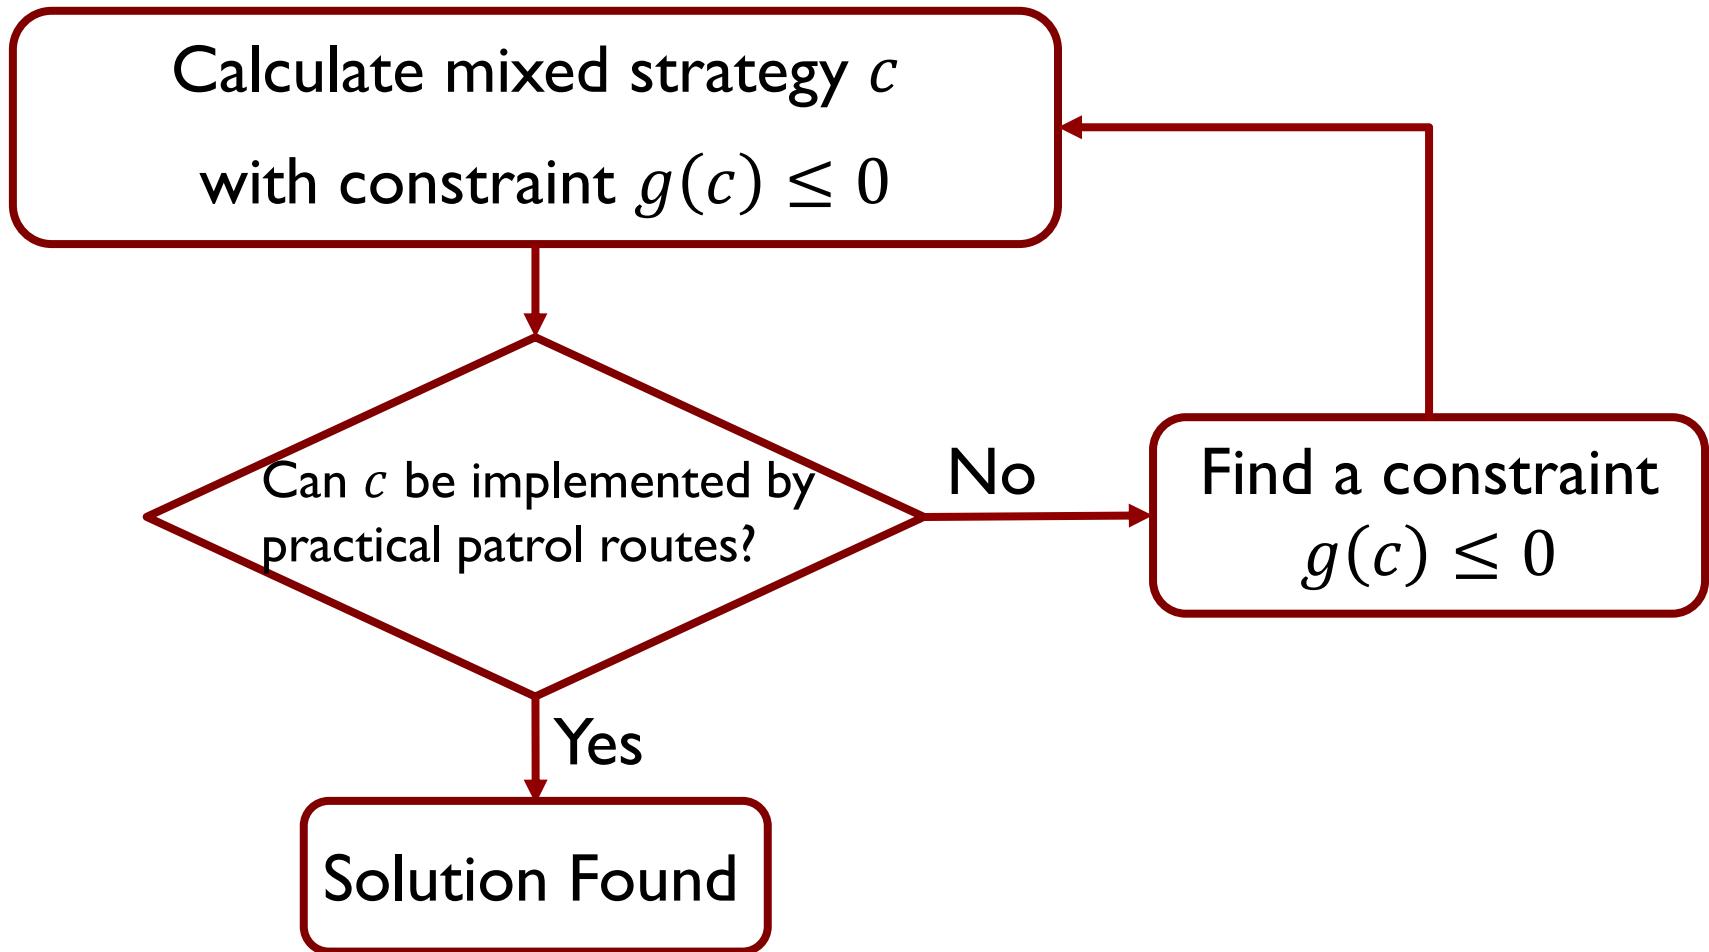


# PAWS Component: Route Planning

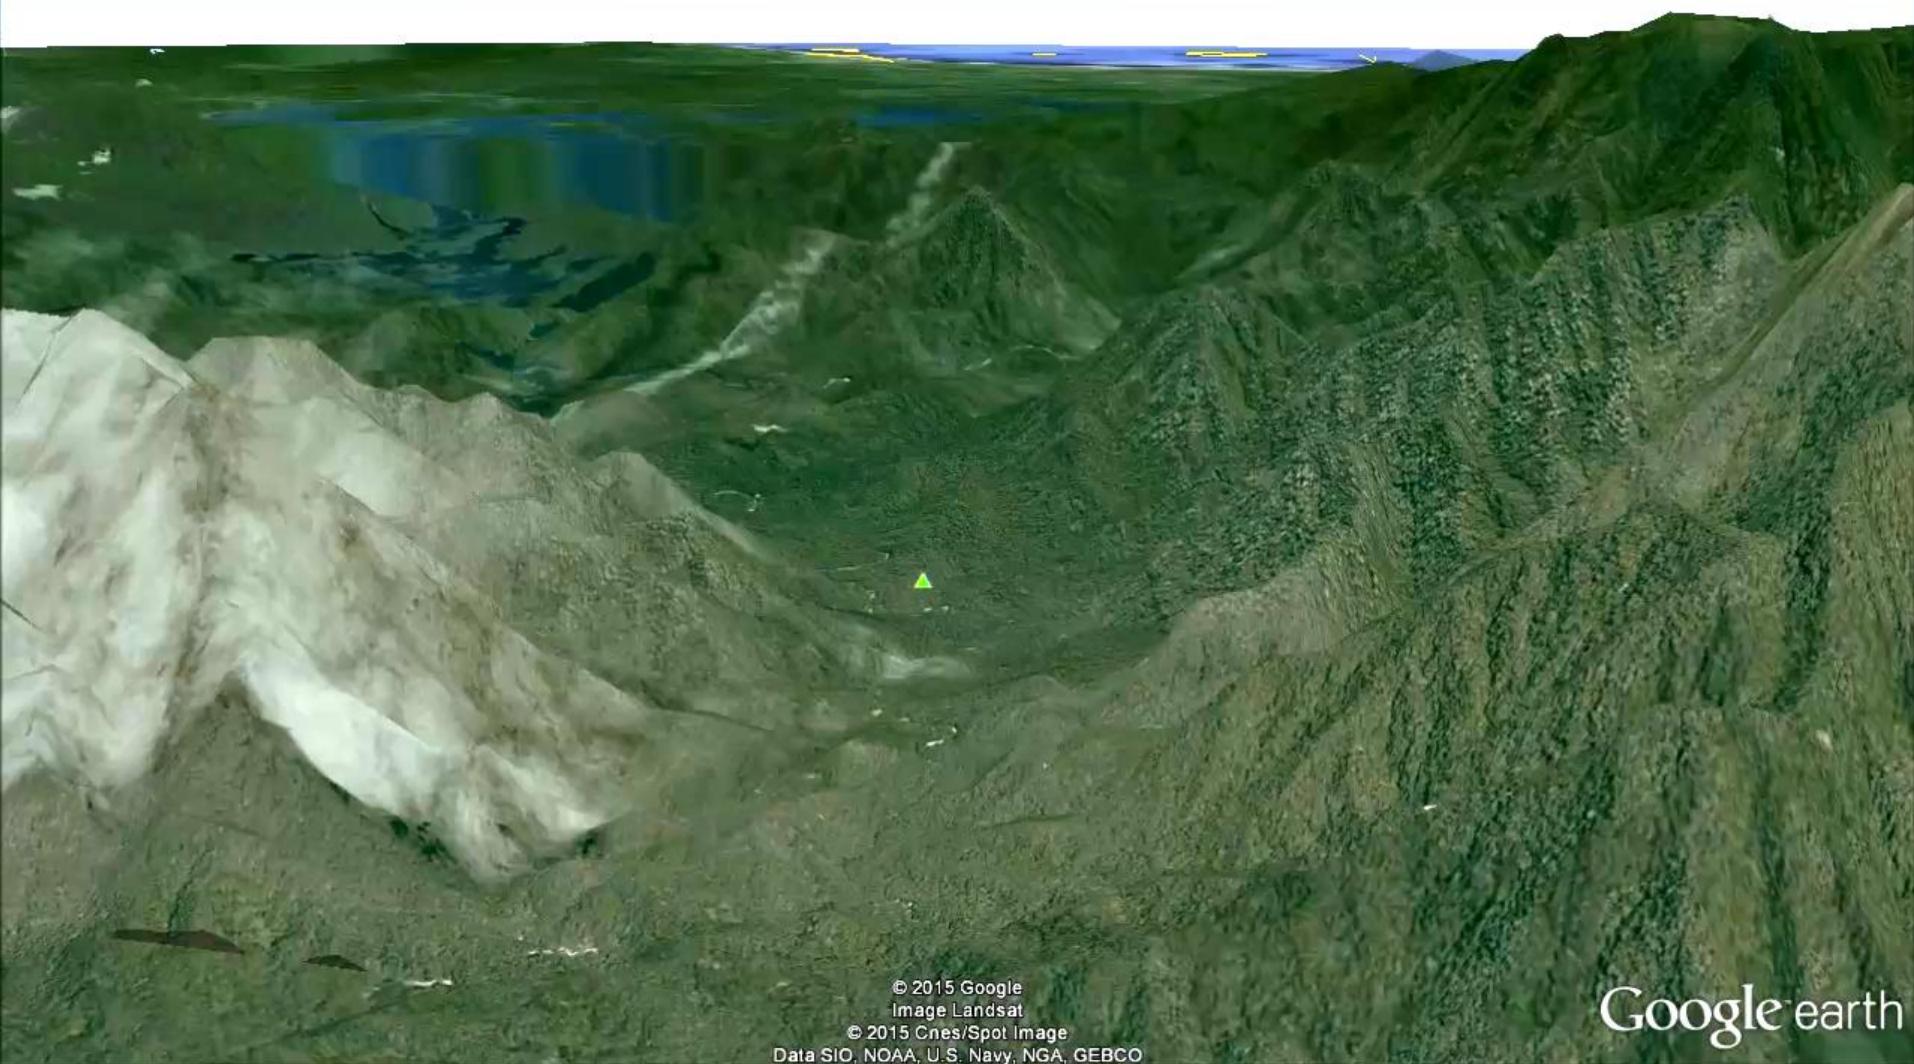
- Hierarchical model: Focus on terrain feature



# PAWS Component: Route Planning

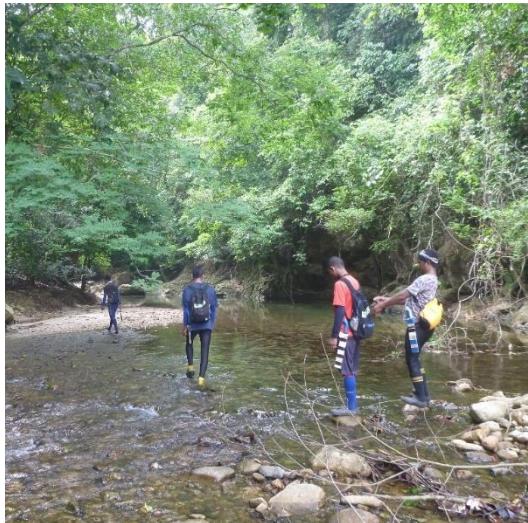


# PAWS Component: Route Planning



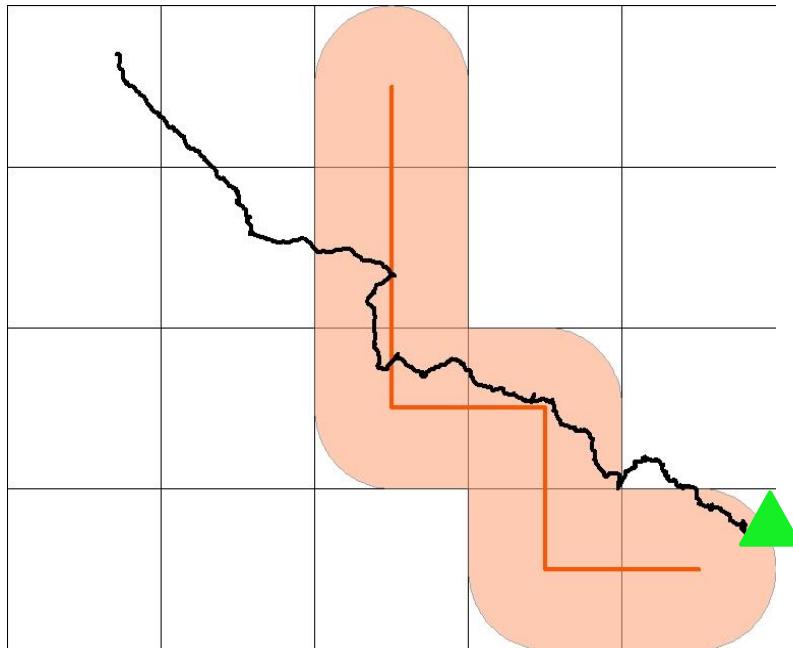
# PAWS Deployment

- ▶ In collaboration with Panthera, Rimba
- ▶ Regular deployment since July 2015 (Malaysia)

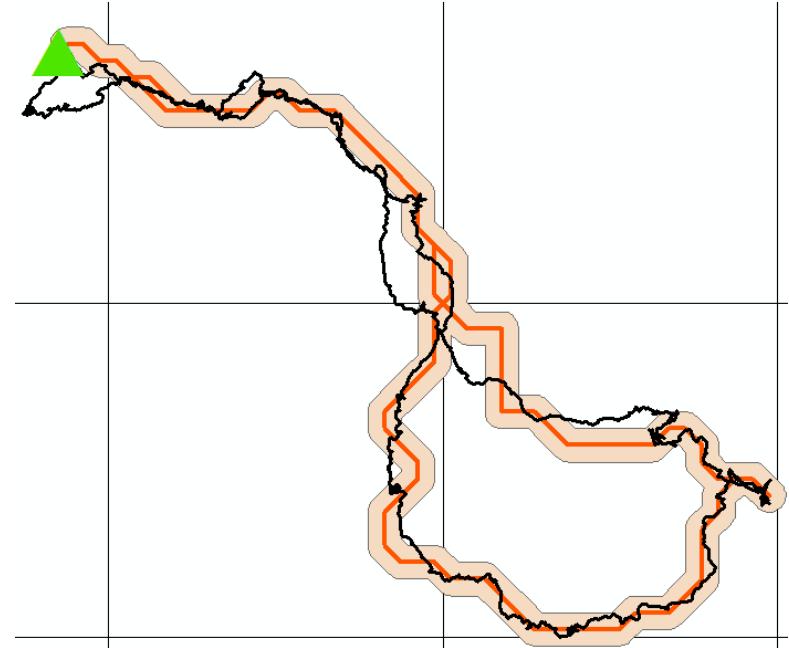


# PAWS Deployment

Grid Based



Route Based



# PAWS Deployment

Animal Footprint



Tree Mark



Tiger Sign

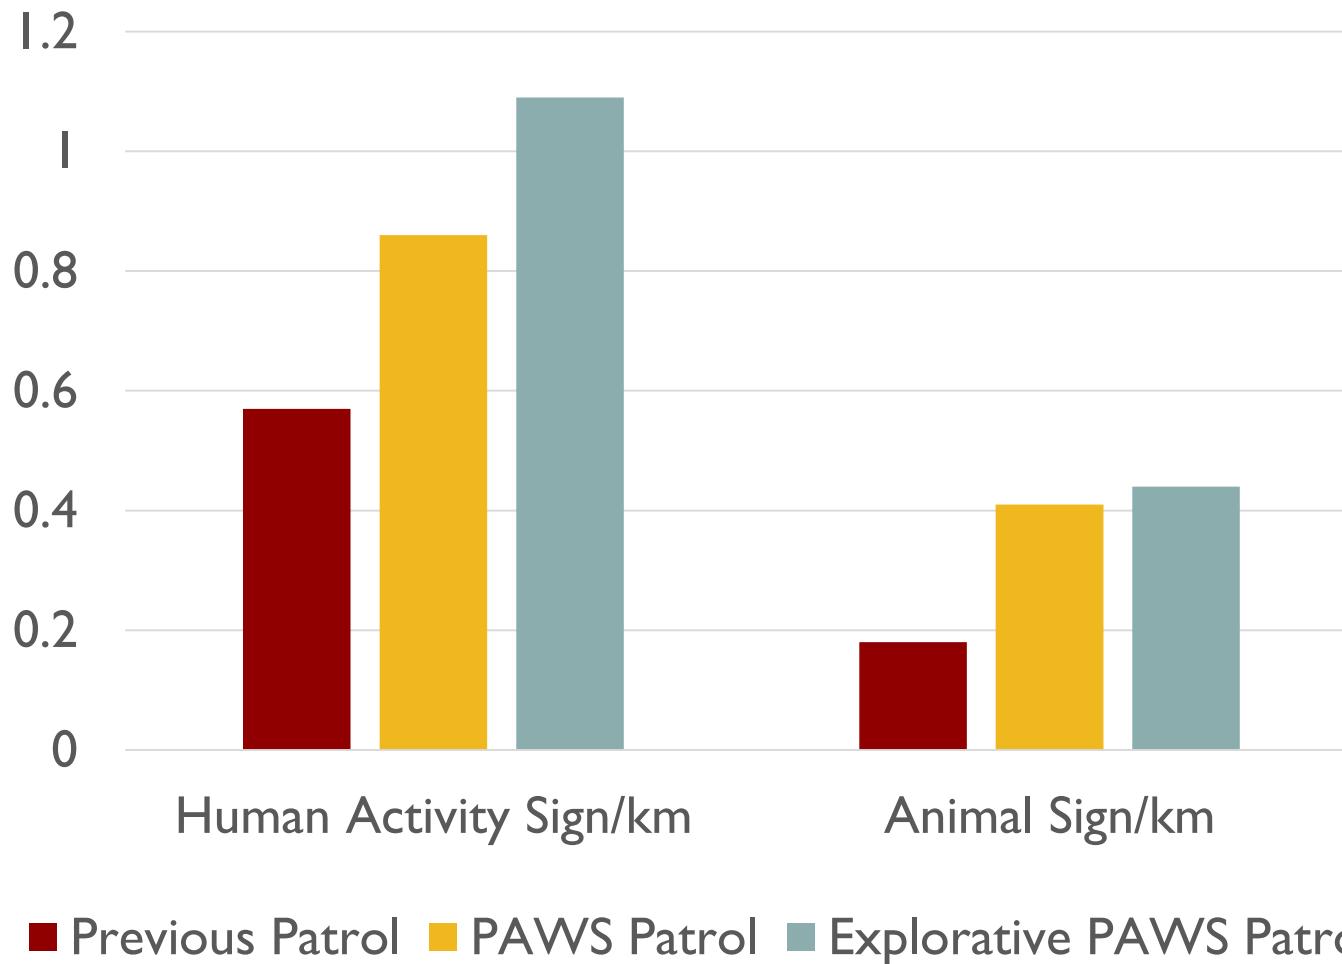


Camping Sign

Lighter



# PAWS Deployment



# PAWS Deployment

- ▶ PAWS is deployed in the field
  - ▶ Saved animals!



# AI and Social Good

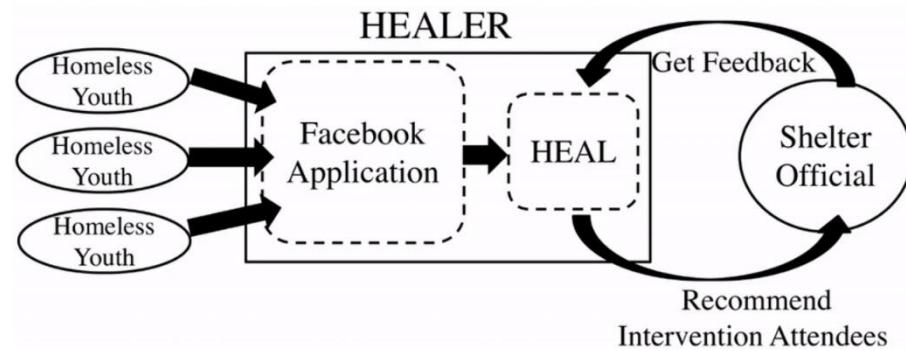
- ▶ AI research that can deliver societal benefits now and in the near future



<http://mashable.com/2015/02/06/hiv-homeless-teens-algorithm/#.k9dRKhxqam>



<https://www.pastemagazine.com/articles/2017/04/a-new-smart-technology-will-help-cities-drastically.html>



# AI and Social Good

► [www.AlandSocialGood.org](http://www.AlandSocialGood.org)



Interest in Artificial Intelligence (AI) has dramatically increased in recent years and AI has been successfully applied to societal challenge problems. It has a great potential to provide tremendous social good in the future.

This website aims to

- Provide information related to AI and Social Good
- Promote research efforts towards AI and Social Good: [Research Expedition](#)
- Highlight social good challenges that are in need of AI research effort: [Social Good Challenges](#)
- Foster collaboration among researchers and practitioners
- Provide resources to researchers and students who are interested in AI and Social Good

# THANK YOU!

Fei Fang  
Harvard University  
Carnegie Mellon University  
[feifang@cmu.edu](mailto:feifang@cmu.edu)